



Small-Space Spectral Sparsification via Bounded-Independence Sampling

DEAN DORON, Ben Gurion University of the Negev, Be'er Sheva, Israel

JACK MURTAGH, Harvard University, Cambridge, MA, USA

SALIL VADHAN and DAVID ZUCKERMAN, University of Texas at Austin, Austin, TX, USA

We give a deterministic, nearly logarithmic-space algorithm for mild spectral sparsification of undirected graphs. Given a weighted, undirected graph G on n vertices described by a binary string of length N , an integer $k \leq \log n$, and an error parameter $\varepsilon > 0$, our algorithm runs in space $\tilde{O}(k \log(N \cdot w_{\max}/w_{\min}))$, where w_{\max} and w_{\min} are the maximum and minimum edge weights in G , and produces a weighted graph H with $\tilde{O}(n^{1+2/k}/\varepsilon^2)$ edges that spectrally approximates G , in the sense of Spielman and Teng, up to an error of ε .

Our algorithm is based on a new bounded-independence analysis of Spielman and Srivastava's effective resistance-based edge sampling algorithm and uses results from recent work on space-bounded Laplacian solvers. In particular, we demonstrate an inherent trade-off (via upper and lower bounds) between the amount of (bounded) independence used in the edge sampling algorithm, denoted by k above, and the resulting sparsity that can be achieved.

CCS Concepts: • **Theory of computation** → **Sparsification and spanners; Pseudorandomness and derandomization**;

Additional Key Words and Phrases: Derandomization, space-bounded computation, graph sparsification

ACM Reference Format:

Dean Doron, Jack Murtagh, Salil Vadhan, and David Zuckerman. 2024. Small-Space Spectral Sparsification via Bounded-Independence Sampling. *ACM Trans. Comput. Theory* 16, 2, Article 7 (March 2024), 32 pages. <https://doi.org/10.1145/3637034>

1 INTRODUCTION

The graph sparsification problem is the following: given a weighted, undirected graph G , compute a graph H that has very few edges but is a close approximation to G for some definition of approximation. In general, graph sparsifiers are useful for developing more efficient graph-theoretic approximation algorithms. Algorithms whose complexity depend on the number of edges in the

D. Doron part of this work was done while at Stanford University.

Dean Doron's research was supported by a Motwani Postdoctoral Fellowship. Jack Murtagh's research was supported by the National Science Foundation (grant no. CCF-1763299). Salil Vadhan's research was supported by the National Science Foundation (grant no. CCF-1763299) and a Simons Investigator Award (grant no. 291814). David Zuckerman's research was supported in part by the National Science Foundation (grant no. CCF-1705028) and a Simons Investigator Award (grant no. 409864).

Authors' addresses: D. Doron, Ben Gurion University of the Negev, David Ben Gurion Blvd 1, Be'er Sheva 84105, Israel; e-mail: deand@bgu.ac.il; J. Murtagh and S. Vadhan, Harvard University, 150 Western Ave, Allston, MA 02134, USA; e-mails: jmurtagh@g.harvard.edu, salil_vadhan@harvard.edu; D. Zuckerman, University of Texas at Austin, 2317 Speedway, Austin, TX 78712, USA; e-mail: diz@cs.utexas.edu.



This work is licensed under a Creative Commons Attribution International 4.0 License.

© 2024 Copyright held by the owner/author(s).

ACM 1942-3454/2024/03-ART7

<https://doi.org/10.1145/3637034>

graph will be more efficient when run on the sparser graph H , and if H approximates G in an appropriate way, the result on H may be a good approximation to the desired result on G . In this work, we present an algorithm that can be implemented deterministically in small space and achieves sparsification in the *spectral sense* of Spielman and Teng [54]. See Section 1.2 for a more formal statement of our main result.

1.1 Background

Motivated by network design and motion planning, Chew [11] studied *graph spanners*, which are sparse versions of graphs that approximately preserve the shortest distance between each pair of vertices. Benczúr and Karger [7] defined *cut sparsifiers* whose notion of approximation is that every cut of H has size within a $(1 \pm \epsilon)$ factor of the size of the corresponding cut in G . They showed that every graph G on n vertices has a cut sparsifier H with $O(n \cdot \log n / \epsilon^2)$ edges and gave a randomized algorithm for computing such cut sparsifiers. Their algorithm runs in nearly linear time (i.e., $\tilde{O}(m)$, where m is the number of edges in G and the $\tilde{O}(\cdot)$ notation hides polylogarithmic factors) and they used it to give a faster algorithm for approximating minimum s - t cuts.

Spielman and Teng introduced *spectral sparsifiers*, which define approximation between the graph and its sparsifier in terms of the quadratic forms of their *Laplacians* [54]. The Laplacian of an undirected graph is the matrix $L = D - A$, where A is the adjacency matrix of the graph and D is the diagonal matrix of vertex degrees (i.e., D_{ii} equals the weighted degree of vertex i). H is said to be an ϵ -*spectral approximation* of G if for all vectors $v \in \mathbb{R}^n$, we have that

$$(1 - \epsilon) \cdot v^\top L v \leq v^\top \tilde{L} v \leq (1 + \epsilon) \cdot v^\top L v,$$

where \tilde{L} and L are the Laplacians of H and G , respectively. Spectral sparsifiers generalize cut sparsifiers, which can be seen by observing that when $v \in \{0, 1\}^n$, v is the characteristic vector of some set of vertices $S \subseteq [n]$ and $v^\top L v$ equals the sum of the weights of the edges cut by S .

Spielman and Teng showed that all graphs have spectral sparsifiers with $O(n \cdot \log^{O(1)} n / \epsilon^2)$ edges and gave a nearly linear time randomized algorithm for computing them with high constant probability. Their spectral sparsifiers were a key ingredient that they used to develop the first nearly linear time algorithm for solving Laplacian systems. These fast Laplacian solvers spawned a flurry of improvements and simplifications [17, 28–30, 32, 33, 36, 46] as well as extensions to directed graphs [14–16] and to the space-bounded setting [1, 19, 43]. Spectral sparsification and the nearly linear time Laplacian solvers that use them have been critical primitives that have enabled the development of faster algorithms for a wide variety of problems, including max flow [12, 18, 26, 35, 39], random generation of spanning trees [27, 42, 51], and other problems in computer science [31, 45].

Spielman and Srivastava [53] gave a spectral sparsification algorithm that both simplified and improved upon the algorithm of Spielman and Teng. They show that randomly sampling edges independently with probabilities proportional to their *effective resistances* produces a good spectral sparsifier with high probability. Viewing a graph as an electrical network, the effective resistance of an edge (a, b) is the potential difference induced between them when a unit of current is injected at a and extracted at b (or vice versa). More formally, the effective resistance of an edge (a, b) in a graph with Laplacian L is

$$R_{ab} = (e_a - e_b)^\top L^+ (e_a - e_b), \quad (1)$$

where e_i denotes the i -th standard basis vector and L^+ denotes the *Moore-Penrose pseudoinverse* of L .¹

¹ L^+ is a matrix with the same kernel as L that acts as an inverse of L on the orthogonal complement of the kernel. See Section 2.2 for a formal definition.

Spielman and Srivastava proved the following theorem.

THEOREM 1.1 (SPECTRAL SPARSIFICATION VIA EFFECTIVE RESISTANCE SAMPLING² [52, 53]). *Let $G = (V, E, w)$ be a weighted graph on n vertices and for each edge $(a, b) \in E$ with weight w_{ab} , define $p_{ab} = \min\{1, 4 \cdot \log n \cdot w_{ab} \cdot R_{ab}/\varepsilon^2\}$, where R_{ab} is the effective resistance of (a, b) as defined in Equation (1). Construct a sparsifier H by sampling edges from G independently such that each edge (a, b) in G is added to H with probability p_{ab} . For edges that get added to H , reweight them with weight w_{ab}/p_{ab} . Let L and \tilde{L} be the Laplacians of G and H , respectively. Then, with high probability,*

- (1) H has $O(n \cdot (\log n)/\varepsilon^2)$ edges and
- (2) \tilde{L} ε -spectrally approximates L .

Furthermore, this procedure can be implemented to run in time $\tilde{O}(\frac{m}{\varepsilon^2} \cdot \log(w_{\max}/w_{\min}))$, where m is the number of edges in G and w_{\max}, w_{\min} are the maximum and minimum edge weights of G , respectively.

The sparsity achieved by the Spielman and Srivastava sparsifiers was improved by Batson et al. [6], who gave a deterministic algorithm for computing ε -spectral sparsifiers with $O(n/\varepsilon^2)$ edges, which is asymptotically optimal. However, their algorithm is less efficient, running in time $O(m \cdot n^3/\varepsilon^2)$. Work on these optimal sparsifiers continued with another slightly faster deterministic algorithm [55] followed by an $O(n^{2+\varepsilon})$ -time randomized algorithm [2] and culminating in the randomized algorithms of Lee and Sun, who achieved almost-linear time [37] and, finally, nearly-linear time [38].

1.2 Our Main Result

In this work, we study the deterministic *space complexity* of computing spectral sparsifiers. We use the standard model of space-bounded computation. We refer the reader to Section 5 (and to [5, Section 4] and [23, Section 5]) for an overview of the model.³

Our main result is a deterministic, nearly-logarithmic space algorithm for computing *mild spectral sparsifiers*, that is, graphs with $O(n^{1+\alpha}/\varepsilon^2)$ edges for any constant $\alpha > 0$.

THEOREM 1.2 (SEE ALSO THEOREM 5.1). *Let G be a connected, weighted, undirected graph on n vertices, $k \in \mathbb{N}$ an independence parameter, and $\varepsilon > 0$ an error parameter. There is a deterministic algorithm that on input G, k , and ε outputs a weighted graph H that is an ε -spectral sparsifier of G and has $O(n^{1+2/k} \cdot (\log n)/\varepsilon^2)$ edges. The algorithm runs in space $O(k \log(N \cdot w) + \log(N \cdot w) \log \log(N \cdot w))$, where $w = w_{\max}/w_{\min}$ is the ratio of the maximum and minimum edge weights in G and N is the length of the input.*

The closest analogue to spectral sparsifiers in the space-bounded derandomization literature is the *derandomized square* of Rozenman and Vadhan [49], a graph operation that produces a sparse approximation to the square of a graph.⁴ The derandomized square was introduced to give an alternative proof to Reingold's celebrated result that Undirected S-T Connectivity can be solved in deterministic logspace [48]. Murtagh et al. [43] showed that the derandomized square actually produces a *spectral sparsifier* of the square of a graph. This was a key observation they used to develop a deterministic, nearly logarithmic space algorithm for solving Laplacian systems. Later, the

²In their original paper [53], they fix the number of edges in the sparsifier in advance, resulting in a slightly different theorem statement and analysis. The version we cite here and what we model our algorithm after was presented later in [52].

³In particular, recall that the input tape is *read only* and the output tape is *write only*. Thus, the space complexity measures the required auxiliary space. To be precise, we measure the space complexity as a function of the input length (in bits) N , rather than the number of vertices n , to also support extreme cases in which $N \gg n$.

⁴The *square* of a graph G is a graph on the same vertex set whose edges correspond to all walks of length 2 in G .

sparsification benefits of the derandomized square were also used in nearly logarithmic space algorithms for deterministically approximating random walk probabilities and for solving Laplacian systems in Eulerian directed graphs [1, 44].

For a d -regular graph G on n vertices, its square G^2 has degree d^2 and the derandomized square computes an ε -spectral approximation to G^2 with degree $O(d/\varepsilon^2)$. On the other hand, applying our sparsification to G^2 results in an ε -spectral approximation with, on average, $O(n^\alpha/\varepsilon^2)$ edges adjacent to each vertex for any constant α , which is independent of d and much sparser when $d = \omega(n^\alpha)$. Also, our algorithm can sparsify any undirected graph, not just squares. Our algorithm does not replace the derandomized square, however, because the derandomized square can be iterated very space efficiently, a property that is used in all of its applications thus far. Nevertheless, given the success of spectral sparsification and Laplacian solvers in the nearly-linear time context and the fruit borne of porting these techniques to the logspace setting, we are hopeful that our spectral sparsifiers will have further applications in derandomization of space-bounded computation.

1.3 Techniques

Our deterministic space-efficient algorithm is modeled after the effective resistance-based sampling algorithm of Spielman and Srivastava (Theorem 1.1). Although the Spielman and Srivastava procedure is randomized and does not achieve optimal sparsity, the known algorithms that do ([2, 6, 37, 38, 55]) are more involved and often sequential in nature. Thus, they do not seem as amenable to small-space implementations.

To derandomize the Spielman–Srivastava algorithm, we follow the standard approach of first reducing the number of random bits used to logarithmic and then enumerating over all random choices of the resulting algorithm. Following [3, 41], a natural way to reduce the number of random bits used is to do the edge sampling only k -wise independently for some $k \ll |E|$ rather than sampling every edge independently from all other edges.

Let k be our bounded-independence parameter. Namely, we are only guaranteed that every subset of k edges is chosen independently (with the right marginals). However, there may be correlations between the choices in tuples of size $k + 1$. It is well known that such a sampling can be performed using fewer random bits. By [53], we know that $k = |E|$ will, with high probability, produce an ε -spectral sparsifier with $O(n \cdot \log n / \varepsilon^2)$ edges in expectation. What about much smaller values of k ? In Section 3, we prove the following:

THEOREM 1.3 (INFORMAL; SEE THEOREM 3.1). *Let G be a connected weighted undirected graph on n vertices with Laplacian L , $k \in \mathbb{N}$ an independence parameter, and $\varepsilon > 0$ an error parameter. Let H be the graph that is the output of Spielman and Srivastava’s sampling-based sparsification algorithm (Theorem 1.1), when the edge sampling is done in a k -wise independent manner, and let \tilde{L} be the Laplacian of H . Then, with high constant probability, \tilde{L} ε -approximates L and H has $O(n^{1+2/k} \cdot (\log n)/\varepsilon^2)$ edges.*

The first thing to observe is that $k = \log n$ gives the same result as in [53]. More importantly, the above shows that the result *interpolates*: Even for a constant k , Theorem 1.3 gives a *mild sparsification* that sparsifies dense graphs to $O(n^{1+\alpha})$ expected edges, where $\alpha > 0$ is an arbitrarily small constant.

We prove Theorem 1.3 by extending the arguments in [52, 53]. For every edge $(a, b) \in E$, we define a random matrix X_{ab} that corresponds to the choice made by the sparsification algorithm, in such a way that $X = \sum_{(a,b) \in E} X_{ab}$ relates to the resulting Laplacian \tilde{L} .⁵ Let Π be the orthogonal

⁵Specifically, $X = L^{+1/2} \tilde{L} L^{+1/2}$, where $L^{+1/2}$ is the square-root of the pseudoinverse of L .

projection onto the image of L . Following [52, 53], we show that \tilde{L} ε -spectrally approximates L (equivalently, that H is an ε -spectral sparsifier for G) with high probability if $X - \Pi$ has bounded moments. Deriving a tail bound that relies on the first k moments alone, we can proceed with the analysis as if the X_{ab} s were *truly independent*. More specifically, we bound $\text{Tr}(\mathbb{E}_X[(X - \Pi)^k])$ using a matrix concentration result due to Chen et al. [10]. We note that matrix tail bounds that rely on the first moments, combined with k -wise independence, was used in previous works, e.g., in [25]. For the complete details, as well as how our argument differs from [52, 53], see Section 3.

Getting a Deterministic Algorithm. Theorem 1.3 readily gives a simple, randomness-efficient algorithm, as k -wise independent sampling of edges only requires $O(k \cdot \log(N \cdot w))$ random bits [3, 24] (see Lemma 2.4). However, more work is needed to obtain a space-efficient deterministic algorithm. First, we need to be able to compute the marginal sampling probabilities, which depend on the effective resistances R_{ab} . Fortunately, the recent work of Murtagh et al. [43] allows us to approximate the effective resistances using only $O(\log(N \cdot w) \log \log(N \cdot w))$ space and we show that the k -wise independent sampling procedure can tolerate the approximation.

Next, to obtain a deterministic algorithm, we can enumerate over all possible random choices of the algorithm in space $O(k \cdot \log(N \cdot w))$ and compute a candidate sparsifier H for each. We are guaranteed that at least one (indeed, most) of the resulting graphs H is a good sparsifier for G but how can we identify which one? To do this, it suffices for us, given Laplacians L and \tilde{L} , to distinguish the case that \tilde{L} is an ε -spectral approximation of L from the case that \tilde{L} is not a 2ε -spectral approximation of L . We reduce that problem to that of approximating the spectral radius of

$$M = \left(\frac{(\tilde{L} - L)L^+}{\varepsilon} \right)^2,$$

where L^+ is the pseudoinverse of L , which can be approximated in nearly logarithmic space by [43]. In fact, it will be sufficient to check whether the trace of a logarithmically high power of M is below a certain threshold to deduce that the spectral radius of M does not exceed 1. In Section 5.2, we show that the latter case implies that \tilde{L} indeed ε -approximates L .

The deterministic, nearly logarithmic space Laplacian solver of [43] only worked for *multi-graphs*, i.e., graphs with integer edge weights. To get our result for arbitrary weighted graphs, we extend the work of [43] and give a deterministic, nearly logarithmic space Laplacian solver for arbitrary undirected weighted graphs. Combining this extension with the k -wise independent analysis of the edge sampling algorithm (Theorem 1.3) and the verification procedure described above lets us prove our main result, Theorem 1.2.

1.4 Lower Bounds for Bounded-Independence Sampling

Having established an upper bound on the amount of independence required for the edge-sampling procedure (Theorem 1.3), a natural goal would be to come up with a corresponding lower bound. Theorem 1.3 tells us that in order to sparsify to $\tilde{O}(n^{1+\alpha})$ expected edges, we can use k -wise independent sampling for $k = 2/\alpha$. Can a substantially smaller choice of k perform just as well? In Section 4, we show that our upper bound of $k = 2/\alpha$ is tight up to a small constant factor.

THEOREM 1.4 (INFORMAL; SEE THEOREM 4.1). *For every small enough $\alpha > 0$, there exist infinitely many connected graphs $G = (V = [n], E)$ with all effective resistances equal that are d -regular with $d = \Omega(n^\alpha)$ and a distribution $\mathcal{D} \sim \{0, 1\}^{|E|}$ that is k -wise independent for $k = \lfloor 4/3\alpha \rfloor$ with marginals $1/2$ that would produce a disconnected graph with high probability.*

Our family of “bad graphs” will be dense graphs having large girth. Namely, given a girth g and an integer $d \geq 3$, we consider graphs $G = (V = [n], E)$ satisfying $d > n^{1/g}$ for some constant $0 < \gamma < 2$ [34]. Getting an infinite family of graphs with γ approaching 2 (and specifically attaining the *Moore bound*), even non-explicitly, has been the subject of extensive study (see [21] and references therein). See also Section 4.1 for a further discussion. Given a sparsification parameter $\alpha > 0$, we set $k \approx \gamma/\alpha$ and take a graph G on n vertices with girth $g = k + 1$ and degree $d > n^{1/g} + 1$.

Our construction of the distribution \mathcal{D} is inspired by Alon and Nussboim [4]: choose a partition of the vertices $V = V_0 \uplus V_1$ uniformly at random, and for every edge $e = (u, v) \in E$, include it in the sample if and only if either $u, v \in V_0$ or $u, v \in V_1$. Clearly, sampling edges according to \mathcal{D} results in a disconnected graph almost surely. However, we show that \mathcal{D} is indeed k -wise independent, relying on the fact that the girth of G is $k + 1$.

To obtain Theorem 1.4, we use the family of graphs given by Lazebnik et al. [34] who obtained $\gamma = 4/3$. Indeed, any improvement in γ would bring our upper bound of $k \approx 2/\alpha$ and lower bound of $k \approx \gamma/\alpha$ closer together.

1.5 Open Problems

An interesting open problem is to achieve improved sparsity, e.g., $O(n \cdot (\log n)/\varepsilon^2)$ matching [53]. Our algorithm would require space $\Omega(\log^2 n)$ to achieve this sparsity, due to setting $k = \Omega(\log n)$. We remark that previous work implies that this can be done in *randomized* logarithmic space. Indeed, Doron et al. [19] gave a randomized algorithm for solving Laplacian systems in logarithmic space (without $\log \log(\cdot)$ factors), which implies that one can approximate effective resistances and, hence, implement the Spielman–Srivastava edge sampling with full independence in randomized logspace. It is also an interesting question as to whether there is a nearly logspace algorithm (even randomized) that produces spectral sparsifiers of optimal sparsity (i.e., $O(n/\varepsilon^2)$ edges).

Finally, there has been recent progress on sparsifying Eulerian digraphs in the nearly-linear time literature [13–16]. Given the recent advance of a nearly-logarithmic space solver for Eulerian Laplacian systems [1], an interesting question is sparsifying Eulerian graphs in small space.

2 PRELIMINARIES

We will work with undirected weighted graphs, $G = (V, E, w)$, where w is a vector of length $|E|$ and each edge $(a, b) \in E$ is associated with a positive weight $w_{ab} > 0$. At times we refer to undirected *multigraphs*, which are weighted graphs where all of the weights are integers. The adjacency matrix of G is a symmetric, real-valued matrix A in which $A_{ij} = w_{ij}$ if $(i, j) \in E$ and $A_{ij} = 0$ otherwise.

For any matrix A , its *spectral norm* $\|A\|$ is $\max_{\|x\|=1} \|Ax\|_2$, which is also the largest singular value of A . For any square matrix A , its *spectral radius*, denoted $\rho(A)$, is the largest absolute value of its eigenvalues. When A is real and symmetric, the spectral norm equals the spectral radius. The spectral norm is sub-multiplicative, i.e., $\|AB\| \leq \|A\| \|B\|$. We denote by A^\top the transpose of A . We denote by $\mathbf{1}$ the all-ones vector, by $\mathbf{0}$ the all-zeros vector, and e_a is the vector with 1 in the a -th coordinate and 0 elsewhere, where e_a 's dimension will be understood from context (i.e., e_a is the a -th standard basis vector).

The *trace* of a matrix $A \in \mathbb{R}^{n \times n}$ is $\text{Tr}(A) = \sum_{i \in [n]} A_{ii}$, which also equals the sum of its eigenvalues. The trace is invariant under cyclic permutations, i.e., $\text{Tr}(AB) = \text{Tr}(BA)$. The expectation of a *random matrix* is the matrix of the coordinate-wise expectations. More formally, if A is a random matrix, then $\mathbb{E}[A] = \widehat{A}$ where $\widehat{A}_{ij} = \mathbb{E}[A_{ij}]$ for all $i, j \in [n]$. The trace and the expectation are both linear functions of a matrix and they commute. That is, for all random matrices A , we have that $\text{Tr}(\mathbb{E}[A]) = \mathbb{E}[\text{Tr}(A)]$ (see, e.g., [47]).

2.1 PSD Matrices and Spectral Approximation

A symmetric matrix $A \in \mathbb{R}^{n \times n}$ is **positive semi-definite (PSD)**, denoted $A \succeq 0$, if for every $x \in \mathbb{R}^n$ it holds that $x^\top A x \geq 0$ or, equivalently, if all its eigenvalues are non-negative. We write $A \succeq B$ if $A - B \succeq 0$.

Definition 2.1. Let A and B be $n \times n$ symmetric PSD matrices. For a real $\varepsilon > 0$, we say that A is an ε -spectral approximation of B , denoted $A \approx_\varepsilon B$, if

$$(1 - \varepsilon)B \leq A \leq (1 + \varepsilon)B.$$

When A and B share an eigenvector basis v_1, \dots, v_n , Definition 2.1 is equivalent to requiring that $(1 - \varepsilon)\mu_i \leq \lambda_i \leq (1 + \varepsilon)\mu_i$, where $\lambda_1, \dots, \lambda_n$ are the eigenvalues of A corresponding to v_1, \dots, v_n and μ_1, \dots, μ_n are the eigenvalues of B corresponding to v_1, \dots, v_n .

2.2 The Moore-Penrose Pseudoinverse

Let A be any linear operator. The *Moore-Penrose pseudoinverse* of A , denoted A^+ , is the unique matrix that satisfies the following:

- (1) $AA^+A = A$,
- (2) $A^+AA^+ = A^+$, and
- (3) both AA^+ and A^+A are Hermitian.

If $A = U\Sigma V^\top$ is the **singular value decomposition (SVD)** of A , the pseudoinverse is given by $A^+ = V\Sigma^+U^\top$, where Σ^+ is the matrix obtained by taking the reciprocal of each nonzero diagonal element of Σ and leaving the zeros intact. When A is a symmetric PSD matrix, the SVD coincides with the eigen-decomposition. Thus, if $\lambda_1, \dots, \lambda_n$ are the eigenvalues of A , then A^+ shares the same eigenvector basis and has eigenvalues $\lambda_1^+, \dots, \lambda_n^+$, where

$$\lambda_i^+ = \begin{cases} 1/\lambda_i & \text{if } \lambda_i \neq 0, \\ 0 & \text{if } \lambda_i = 0. \end{cases}$$

Also, note that if A is real then A^+ is real valued as well.

A *square root* of a matrix A is any matrix X that satisfies $X^2 = A$. When A is symmetric and PSD, it has a unique symmetric PSD square root, which we write as $A^{1/2}$. If $A = U\Sigma U^\top$ is the eigen-decomposition of A , then $A^{1/2} = U\sqrt{\Sigma}U^\top$ where $\sqrt{\Sigma}$ is obtained by taking the square root of each diagonal element of Σ . We denote by $A^{+/2}$ the matrix $(A^+)^{1/2} = (A^{1/2})^+$.

2.3 The Graph Laplacian and Effective Resistance

Given a graph G on n vertices with an adjacency matrix A and degree matrix D (i.e., D is a diagonal matrix where $D_{ii} = \sum_{j=1}^n A_{ij}$ equals the weighted degree of vertex i in G), the *Laplacian* of G is the matrix

$$L = D - A.$$

For every undirected weighted graph $G = (V, E, w)$, its Laplacian L is symmetric and PSD, with smallest eigenvalue 0. The zero eigenvalue has multiplicity one if and only if G is connected. In this case, $\ker(L_G) = \text{span}(\{1\})$. For every edge $(a, b) \in E$, define the *edge Laplacian* of (a, b) to be

$$L_{ab} = (e_a - e_b)(e_a - e_b)^\top = (e_b - e_a)(e_b - e_a)^\top.$$

Note that $L = \sum_{(a,b) \in E} w_{ab} \cdot L_{ab}$.

It is often helpful to associate G with an electric circuit, where an edge $(a, b) \in E$ corresponds to a resistor of resistance $1/w_{ab}$. For each pair of vertices a and b , the *effective resistance* between

them, denoted by R_{ab} , is the energy of the electrical flow that sends one unit of current from a to b . The effective resistance can be calculated using the pseudoinverse of the Laplacian:

$$R_{ab} = (e_a - e_b)^\top L^+ (e_a - e_b).$$

See [9] for more information on Laplacians and viewing graphs as electrical networks. A useful fact about effective resistances is Foster's Theorem:

THEOREM 2.2 ([22]). *For every undirected connected weighted graph $G = (V, E, w)$ on n vertices, it holds that*

$$\sum_{(a,b) \in E} w_{ab} \cdot R_{ab} = n - 1.$$

2.4 Bounded-Independence Sampling

Given a probability vector $p \in [0, 1]^m$, let $\text{Bernoulli}(p)$ denote the distribution X over $\{0, 1\}^m$, where the bits are independent and for each $i \in [m]$, $\mathbb{E}[X_i] = p_i$. For a set $I \subseteq [m]$ and a string $z \in \{0, 1\}^m$, we let $z|_I \in \{0, 1\}^{|I|}$ be the restriction of z to the indices in I .

Definition 2.3. We say that a distribution $X \sim \{0, 1\}^m$ is k -wise independent with marginals $p \in [0, 1]^m$ if for every set $I \subseteq [m]$ with $|I| \leq k$, it holds that $X|_I = \text{Bernoulli}(p|_I)$. We refer to X as a k -wise independent sample space with marginals p .

Consider $G = (V, E, w)$, with $|E| = m$. Throughout, when we say *sampling edges in a k -wise independent manner*, we refer to the process of picking an element $x \in \{0, 1\}^m$ from a k -wise independent sample space uniformly at random and taking those edges $e \in E$ for which $x_e = 1$.

For $p \in [0, 1]^m$ and a positive integer t , we define $\lfloor p \rfloor_t$ to be the vector p' obtained by truncating every element of p after t bits. Thus, for each $i \in [m]$, $p'_i = 2^{-t} \lfloor 2^t p_i \rfloor$; therefore, $|p_i - p'_i| \leq 2^{-t}$. The following lemma states that we can construct small k -wise independent sample spaces with any specified marginals.

LEMMA 2.4 (FOLLOWING [3, 24]). *For every $m, k, t \in \mathbb{N}$ and $p \in [0, 1]^m$, there exists an explicit k -wise independent distribution $X \sim \{0, 1\}^m$ with marginals $\lfloor p \rfloor_t$ that can be sampled with $r = O(k \cdot \max\{t, \log m\})$ truly random bits. Furthermore, given $\rho \in \{0, 1\}^r$, the element $x \in \text{Supp}(X)$ corresponding to the random bits ρ can be computed in $O(k \cdot \max\{t, \log m\})$ space.*

2.5 Auxiliary Claims

We will need the following claims, whose proofs we will defer to Appendix A.

CLAIM 2.5. *Let A, B, C be $n \times n$ symmetric PSD matrices and suppose that $B \leq C$. Then,*

$$\|A + B\| \leq \|A + C\|.$$

CLAIM 2.6. *Let $G = (V, E, w)$ be an undirected weighted graph on n vertices with Laplacian L . Fix $(a, b) \in E$ and recall that $L_{ab} = (e_a - e_b)(e_a - e_b)^\top$. Then,*

$$\left\| L^{+1/2} L_{ab} L^{+1/2} \right\| = R_{ab}.$$

CLAIM 2.7. *Let $G = (V, E, w)$ be an undirected connected weighted graph on n vertices with Laplacian L . Let J be the $n \times n$ matrix with $1/n$ in every entry and define $\Pi = I - J$ (i.e., Π is the projection onto $\text{span}(\mathbf{1})^\perp = \text{Im}(L)$). Then, we have that*

$$\Pi = LL^+ = L^+L = L^{+1/2}LL^{+1/2}.$$

CLAIM 2.8. *Let A, B, C be symmetric $n \times n$ matrices and suppose that A and B are PSD. Then, the following hold:*

Sparsify($G = (V, E, w), \{\tilde{R}_{ab}\}_{(a,b) \in E}, k, \varepsilon, \delta$)

- (1) Initialize H to be the empty graph on $n = |V(G)|$ vertices.
- (2) Set $s \leftarrow \frac{18\varepsilon \log n}{\varepsilon^2} \cdot \left(\frac{n}{\delta}\right)^{2/k}$.
- (3) For every edge $(a, b) \in E$, set $p_{ab} \leftarrow \min \left\{1, w_{ab} \cdot \tilde{R}_{ab} \cdot s\right\}$
- (4) For every edge $(a, b) \in E$, add (a, b) to H with weight w_{ab}/p_{ab} with probability p_{ab} . Do this sampling in a k -wise independent manner, following Lemma 2.4.
- (5) Return H .

Fig. 1. Computing a spectral sparsifier via bounded independence sampling.

- (1) $A \approx_\varepsilon B \implies C^\top AC \approx_\varepsilon C^\top BC$
- (2) If $\ker(C) \subseteq \ker(A) = \ker(B)$, then $A \approx_\varepsilon B \iff C^\top AC \approx_\varepsilon C^\top BC$

The proof of the following claim can be found in [54].

CLAIM 2.9 ([54]). *Let G be an undirected, weighted graph on n vertices with Laplacian L and minimum weight w_{\min} . Then, the smallest nonzero eigenvalue of L is at least $\min\left\{\frac{8w_{\min}}{n^2}, \frac{w_{\min}}{n}\right\}$.*

3 SPARSIFICATION VIA BOUNDED-INDEPENDENCE SAMPLING

In Section 1, we briefly introduced the Spielman–Srivastava sparsification algorithm [53] based on (truly) independent edge sampling, with probabilities proportional to the effective resistances of the edges. In this section, we explore the trade-off between the amount of independence used in the edge sampling process and the resulting sparsity that can be achieved.

In particular, we analyze the algorithm Sparsify (see Figure 1). The algorithm gets as input an undirected, weighted, dense graph $G = (V, E, w)$ on n vertices, approximate effective resistances \tilde{R}_{ab} for each edge $(a, b) \in E$, a bounded independence parameter $k \leq \log n$, a desired approximation error $\varepsilon > 0$, and a parameter $\delta > 0$ governing the success probability, and outputs a sparser graph H whose Laplacian ε -spectral approximates the Laplacian of G with probability at least $1 - 2\delta$.

First, we will analyze Sparsify for the case in which the effective resistances are given exactly, i.e., $\tilde{R}_{ab} = R_{ab}$ for all $(a, b) \in E$. Then, in Section 3.2, we will analyze the more general case in which we are given approximations to the effective resistances. This latter case is useful algorithmically because more efficient algorithms are known for estimating effective resistances rather than for computing them exactly, both in the time-bounded and space-bounded settings [43, 53].

3.1 Sparsification With Exact Effective Resistances

In this section, we prove the following theorem about Sparsify.

THEOREM 3.1 (SPECTRAL SPARSIFICATION VIA BOUNDED INDEPENDENCE). *Let $G = (V, E, w)$ be an undirected connected weighted graph on n vertices with Laplacian L and effective resistances $R = \{R_{ab}\}_{(a,b) \in E}$. Let $0 < \varepsilon < 1$, $0 < \delta < 1/2$ and let $k \leq \log n$ be an even integer. Let H be the output of Sparsify($G, R, k, \varepsilon, \delta$) and let \tilde{L} be its Laplacian. Then, with probability at least $1 - 2\delta$, we have that*

- (1) $\tilde{L} \approx_\varepsilon L$ and
- (2) H has $O\left(\frac{1}{\delta^{1+2/k}} \cdot \frac{\log n}{\varepsilon^2} \cdot n^{1+\frac{2}{k}}\right)$ edges.

Spielman and Srivastava showed that by using truly independent sampling (i.e., $k = |E|$) in Sparsify, one can compute an ε -spectral sparsification of G with $O(n \cdot \log n / \varepsilon^2)$ edges, with high constant probability [53]. One immediate consequence of Theorem 3.1 is that $\log n$ -wise independent sampling suffices to match the sparsity that truly independent sampling achieves. Another

consequence of Theorem 3.1 is that for any constant $0 < \alpha < 1$ and any constant $\gamma < \alpha/2$, for $k \approx 2/(\alpha - 2\gamma)$, k -wise independent sampling achieves a spectral sparsifier with error $\varepsilon = n^{-\gamma}$ and $O(n^{1+\alpha})$ expected edges, with high constant probability.

The proof of Theorem 3.1 is modeled after Spielman and Srivastava's argument [53]. One difference is that the sparsification algorithm in [53] fixes the number of edges to be sampled in advance rather than having the number of edges be a random variable. They then prove spectral approximation by reducing the problem to a question about concentration of random matrices, which they resolve with a matrix Chernoff bound due to Rudelson and Vershynin [50]. We follow a variant of this argument for the case in which the number of edges in the sparsifier is random and use a matrix concentration bound of Chen et al. [10]. This variant, for truly independent sampling, has appeared before in [52]. Our argument deviates in the proof of Lemma 3.7 to address the fact that we only use k -wise independent sampling.

We start by showing that the sparsity guarantee in Theorem 3.1 indeed holds. Since the inclusion or exclusion of each edge in the sparsifier is a Bernoulli random variable, we can write the expected number of edges in it as

$$\begin{aligned} \sum_{(a,b) \in E} p_{ab} &\leq \sum_{(a,b) \in E} w_{ab} \cdot R_{ab} \cdot s \\ &= (n-1) \cdot s \\ &= O\left(\frac{\log n}{\delta^{2/k} \varepsilon^2} \cdot n^{\frac{2}{k}}\right) \cdot n, \end{aligned}$$

where the second line follows from Claim 2.2. By Markov's inequality, we can conclude the following.

CLAIM 3.2. *Item (2) of Theorem 3.1 holds with probability at least $1 - \delta$.*

We prove item (1) of Theorem 3.1 by the following sequence of lemmas. Throughout, we let $G = (V, E, w)$, L , \tilde{L} , ε , and k as in Theorem 3.1 and s , R_{ab} , and p_{ab} as in Sparsify. Let $\Pi = I - J$ be the orthogonal projection onto $\text{Im}(L)$, as in Claim 2.7. For each $(a, b) \in E$, we define the random matrix

$$X_{ab} = \begin{cases} \frac{w_{ab}}{p_{ab}} \cdot L^{+/2} L_{ab} L^{+/2} & \text{if we choose to include edge } (a, b) \text{ in Step 4 of Sparsify} \\ 0 & \text{otherwise.} \end{cases}$$

Thus, $\mathbb{E}[X_{ab}] = w_{ab} \cdot L^{+/2} L_{ab} L^{+/2}$ and the X_{ab} s are k -wise independent. That is,

$$(X_{e_1}, \dots, X_{e_k}) \equiv X_{e_1} \times \dots \times X_{e_k}$$

for every $\{e_1, \dots, e_k\} \subseteq E$. Let $X = \sum_{(a,b) \in E} X_{ab}$. Also, recall that $L_{ab} = (e_a - e_b)(e_a - e_b)^\top$.

LEMMA 3.3. *Fix all the random choices in Sparsify. Then, $\tilde{L} \approx_\varepsilon L$ if and only if $X \approx_\varepsilon \Pi$.*

PROOF. By definition, $\tilde{L} \approx_\varepsilon L$ if and only if

$$(1 - \varepsilon)L \leq \tilde{L} \leq (1 + \varepsilon)L.$$

Multiplying on both sides by $L^{+/2}$ and applying Claims 2.7 and 2.8, we get that this is equivalent to

$$(1 - \varepsilon)\Pi \leq L^{+/2} \tilde{L} L^{+/2} \leq (1 + \varepsilon)\Pi.$$

For each $(a, b) \in E$, let Y_{ab} be the random variable

$$Y_{ab} = \begin{cases} \frac{w_{ab}}{p_{ab}} \cdot L_{ab} & \text{if we choose to include edge } (a, b) \text{ in Step 4 of Sparsify} \\ 0 & \text{otherwise.} \end{cases}$$

Note that $\sum_{(a,b) \in E} Y_{ab} = \tilde{L}$. Thus, we have that

$$\begin{aligned} X &= \sum_{(a,b) \in E} X_{ab} \\ &= L^{+/2} \left(\sum_{(a,b) \in E} Y_{ab} \right) L^{+/2} \\ &= L^{+/2} \tilde{L} L^{+/2}. \end{aligned}$$

It follows that $\tilde{L} \approx_\varepsilon L$ if and only if $X \approx_\varepsilon \Pi$. \square

LEMMA 3.4. *Fix all the random choices in Sparsify and assume that k is even. Then, $X \approx_\varepsilon \Pi$ if $(X - \Pi)^k \leq (\varepsilon \cdot \Pi)^k$.*

PROOF. First, we observe that X and Π share a common eigenbasis. $\mathbf{1}$ is in the kernel of both $\Pi = I - J$ and X . Let v_2, \dots, v_n be orthogonal eigenvectors of X in $\text{span}(\{\mathbf{1}\})^\perp$ with eigenvalues $\lambda_2, \dots, \lambda_n$, respectively. These are all also eigenvectors of Π since $\text{span}(\{\mathbf{1}\})^\perp$ is an eigenspace of Π of eigenvalue 1. Assume that $(X - \Pi)^k \leq (\varepsilon \cdot \Pi)^k$. Since v_i is an eigenvector of both X and Π , we have that $(\lambda_i - 1)^k \leq \varepsilon^k$. Since k is even, it follows that $|\lambda_i - 1| \leq \varepsilon$, i.e., $1 - \varepsilon \leq \lambda_i \leq 1 + \varepsilon$. Since this holds for all eigenvalues $\lambda_2, \dots, \lambda_n$ of X and the corresponding eigenvalues of Π are 1, we conclude that $X \approx_\varepsilon \Pi$. \square

LEMMA 3.5. *Fix all the random choices in Sparsify and assume that k is even. Then, $(X - \Pi)^k \leq (\varepsilon \cdot \Pi)^k$ if $\text{Tr}((X - \Pi)^k) \leq \varepsilon^k$.*

PROOF. Since k is even, $(X - \Pi)^k$ has non-negative eigenvalues. If $\text{Tr}((X - \Pi)^k) \leq \varepsilon^k$, then the sum of the eigenvalues of $(X - \Pi)^k$ is at most ε^k . Hence, the largest eigenvalue of $(X - \Pi)^k$ is at most ε^k . Since $\ker(X) \supseteq \ker(\Pi)$ and all non-zero eigenvalues of $(\varepsilon \cdot \Pi)^k$ equal ε^k , it follows that $(X - \Pi)^k \leq (\varepsilon \cdot \Pi)^k$. \square

LEMMA 3.6. *Assuming that k is even, it holds that $\Pr_X[\text{Tr}((X - \Pi)^k) > \varepsilon^k] \leq \frac{1}{\varepsilon^k} \text{Tr}(\mathbb{E}_X[(X - \Pi)^k])$.*

PROOF. Since $\text{Tr}((X - \Pi)^k)$ is non-negative (due to k being even), Markov's inequality gives

$$\Pr_X[\text{Tr}((X - \Pi)^k) > \varepsilon^k] \leq \frac{1}{\varepsilon^k} \cdot \mathbb{E}_X[\text{Tr}((X - \Pi)^k)].$$

Note that the trace and the expectation commute completes the proof. \square

LEMMA 3.7. *It holds that $\text{Tr}(\mathbb{E}_X[(X - \Pi)^k]) \leq n \cdot (\frac{18e \log n}{s})^{k/2}$.*

To prove Lemma 3.7, we will use the following theorem of Chen, Gittens, and Tropp.

THEOREM 3.8 ([10]). *Let W_1, \dots, W_m be independent, random, symmetric $n \times n$ matrices. Fix $k \geq 2$ and let $r = \max\{k, 2 \log n\}$. Then,*

$$\left(\mathbb{E} \left[\left\| \sum_{i \in [m]} W_i \right\|^k \right] \right)^{1/k} \leq \sqrt{e \cdot r} \cdot \left\| \sum_{i \in [m]} \mathbb{E}[W_i^2] \right\|^{1/2} + 2e \cdot r \cdot \left(\mathbb{E} \left[\max_{i \in [m]} \|W_i\|^k \right] \right)^{1/k}.$$

PROOF OF LEMMA 3.7. Define $Z_{ab} = X_{ab} - w_{ab} \cdot L^{+1/2} L_{ab} L^{+1/2}$ and let $Z = \sum_{(a,b) \in E} Z_{ab} = X - \Pi$. Our goal is to bound $\mathbb{E}[Z^k]$. To do this, we define

$$\left\{ \widehat{Z}_{ab} \right\}_{(a,b) \in E}$$

to be identically distributed to the Z_{ab} s, except that the \widehat{Z}_{ab} random variables are truly independent instead of only k -wise. More specifically, if we let \widehat{X}_{ab} be defined the same way as X_{ab} but this time we sample the edges in Step 4 of Sparsify truly independently (with marginals p_{ab}), then $\widehat{Z}_{ab} = \widehat{X}_{ab} - w_{ab} \cdot L^{+1/2} L_{ab} L^{+1/2}$. Analogously, let $\widehat{Z} = \sum_{(a,b) \in E} \widehat{Z}_{ab}$.

The key point to note is that both $\mathbb{E}[\widehat{Z}^k]$ and $\mathbb{E}[Z^k]$ can each be written as a sum of products of at most k random variables. As the Z_{ab} s are k -wise independent, we have the following.

CLAIM 3.9. $\mathbb{E}[\widehat{Z}^k] = \mathbb{E}[Z^k]$.

Towards bounding $\text{Tr}(\mathbb{E}[\widehat{Z}^k]) = \mathbb{E}[\text{Tr}(\widehat{Z}^k)]$, we first bound $\mathbb{E}[\|\widehat{Z}^k\|]$. Then, we use the fact that for all symmetric $n \times n$ matrices M , we have that $\text{Tr}(M) \leq n \cdot \|M\|$. Thus,

$$\text{Tr}(\mathbb{E}[\widehat{Z}^k]) \leq n \cdot \mathbb{E}[\|\widehat{Z}^k\|] \leq n \cdot \mathbb{E}[\|\widehat{Z}\|^k],$$

where the latter inequality is by the submultiplicity of the spectral norm. Since $\widehat{Z} = \sum_{(a,b) \in E} \widehat{Z}_{ab}$, we can bound the right-hand side by applying Theorem 3.8 to the \widehat{Z}_{ab} s. To bound the two terms on the right-hand side, we make use of the following two claims.

CLAIM 3.10. For every $(a, b) \in E$ and every matrix in the support of \widehat{Z}_{ab} , it holds that $\|\widehat{Z}_{ab}\| \leq \frac{1}{s}$.

PROOF. Observe that if $p_{ab} = 1$, then $\widehat{Z}_{ab} = 0$. If $p_{ab} < 1$,

$$\begin{aligned} \|\widehat{Z}_{ab}\| &\leq \max \left\{ \left\| \left(\frac{1}{p_{ab}} - 1 \right) \cdot w_{ab} \cdot L^{+1/2} L_{ab} L^{+1/2} \right\|, \left\| -w_{ab} \cdot L^{+1/2} L_{ab} L^{+1/2} \right\| \right\} \\ &\leq \frac{1}{p_{ab}} \cdot \left\| w_{ab} \cdot L^{+1/2} L_{ab} L^{+1/2} \right\| \\ &= \frac{w_{ab} \cdot R_{ab}}{p_{ab}} \quad (\text{Claim 2.6}) \\ &= \frac{1}{s}. \end{aligned}$$

□

CLAIM 3.11. For every $(a, b) \in E$, it holds that $\mathbb{E}[\widehat{Z}_{ab}^2] \leq \frac{w_{ab}}{s} \cdot L^{+1/2} L_{ab} L^{+1/2}$.

PROOF. As $\mathbb{E}[\widehat{Z}_{ab}] = 0$, we can write

$$\begin{aligned} \mathbb{E}[\widehat{Z}_{ab}^2] &= \text{Var}[X_{ab}] \\ &= \frac{1}{p_{ab}^2} \cdot \text{Var}[\text{Ber}(p_{ab})] \cdot \left(w_{ab} \cdot L^{+1/2} L_{ab} L^{+1/2} \right)^2 \\ &= \left(\frac{1}{p_{ab}} - 1 \right) \cdot \left(w_{ab} \cdot L^{+1/2} L_{ab} L^{+1/2} \right)^2 \\ &= \left(\frac{1}{p_{ab}} - 1 \right) \cdot w_{ab}^2 \cdot L^{+1/2} (e_a - e_b)(e_a - e_b)^\top L^+ (e_a - e_b)(e_a - e_b)^\top L^{+1/2} \end{aligned}$$

$$\begin{aligned}
&= \left(\frac{1}{p_{ab}} - 1 \right) \cdot w_{ab}^2 \cdot L^{+1/2} (e_a - e_b) R_{ab} (e_a - e_b)^\top L^{+1/2} \\
&= R_{ab} \cdot w_{ab}^2 \cdot \left(\frac{1}{p_{ab}} - 1 \right) \cdot L^{+1/2} L_{ab} L^{+1/2}.
\end{aligned}$$

Note that if $p_{ab} = 1$, then the above expectation is 0. If $p_{ab} < 1$, then

$$R_{ab} \cdot w_{ab}^2 \cdot \left(\frac{1}{p_{ab}} - 1 \right) \cdot L^{+1/2} L_{ab} L^{+1/2} \leq \frac{w_{ab}}{s} \cdot L^{+1/2} L_{ab} L^{+1/2}.$$

□

Now we can bound the first term on the right-hand side of Theorem 3.8. Together, Claims 3.11 and 2.5 give:

$$\left\| \sum_{(a,b) \in E} \mathbb{E}[\widehat{Z}_{ab}^2] \right\| \leq \left\| \frac{1}{s} \cdot \sum_{(a,b) \in E} w_{ab} \cdot L^{+1/2} L_{ab} L^{+1/2} \right\|.$$

Now, recall that $\sum_{(a,b) \in E} w_{ab} \cdot L^{+1/2} L_{ab} L^{+1/2} = L^{+1/2} L L^{+1/2} = \Pi$ and $\|\Pi\| = 1$. Thus,

$$\left\| \sum_{(a,b) \in E} \mathbb{E}[\widehat{Z}_{ab}^2] \right\|^{1/2} \leq \frac{1}{\sqrt{s}}.$$

To bound the second term of Theorem 3.8, we apply Claim 3.10 to get

$$\mathbb{E} \left[\max_{(a,b) \in E} \|\widehat{Z}_{ab}\|^k \right]^{1/k} \leq \frac{1}{s}.$$

Set $r = \max \{k, 2 \log n\} = 2 \log n$. Combining the bounds on the two terms and applying Theorem 3.8 gives

$$\begin{aligned}
\left(\mathbb{E} \left[\|\widehat{Z}\|^k \right] \right)^{1/k} &\leq \sqrt{\frac{e \cdot r}{s}} + \frac{2e \cdot r}{s} \\
&\leq 3 \cdot \sqrt{\frac{2e \log n}{s}},
\end{aligned}$$

when $s > e \cdot r$. Raising both sides to the k -th power and using the sub-multiplicativity of the spectral norm, we get

$$\mathbb{E} \left[\|\widehat{Z}^k\| \right] \leq \mathbb{E} \left[\|\widehat{Z}\|^k \right] \leq \left(\frac{18e \log n}{s} \right)^{k/2}.$$

For all symmetric $n \times n$ matrices M , we have that $\text{Tr}(M) \leq n \cdot \|M\|$. Thus, by the monotonicity of expectation, we get that

$$\text{Tr}(\mathbb{E}[\widehat{Z}^k]) = \mathbb{E} \left[\text{Tr}(\widehat{Z}^k) \right] \leq n \cdot \left(\frac{18e \log n}{s} \right)^{k/2}.$$

By Claim 3.9,

$$\mathbb{E} \left[\text{Tr}(Z^k) \right] = \text{Tr}(\mathbb{E}[Z^k]) \leq n \cdot \left(\frac{18e \log n}{s} \right)^{k/2},$$

and, thus,

$$\frac{1}{\varepsilon^k} \cdot \text{Tr}(\mathbb{E}[(X - \Pi)^k]) \leq \frac{n}{\varepsilon^k} \cdot \left(\frac{18e \log n}{s} \right)^{k/2}.$$

□

Now we can prove the main theorem of this section.

PROOF OF THEOREM 3.1. From Lemma 3.3, Lemma 3.4, Lemma 3.5, and Lemma 3.6, we have that $\tilde{L} \approx_\varepsilon L$ except with probability at most

$$\frac{1}{\varepsilon^k} \cdot \text{Tr}(\mathbb{E}[(X - \Pi)^k]).$$

By Lemma 3.7, we have that

$$\frac{1}{\varepsilon^k} \cdot \text{Tr}(\mathbb{E}[(X - \Pi)^k]) \leq \frac{n}{\varepsilon^k} \cdot \left(\frac{18e \log n}{s} \right)^{k/2}.$$

The above is upper bounded by δ whenever

$$s \geq \frac{18e \log n \cdot n^{2/k}}{\delta^{2/k} \varepsilon^2},$$

which is how we set s in Sparsify. Combining this with Claim 3.2, the theorem follows by the union bound. \square

3.2 Sparsification With Approximate Effective Resistances

Spielman and Srivastava showed that the original version of spectral sparsification through effective resistance sampling (with fully independent sampling and fixing the number of edges in advance) is robust to small changes in the sampling probabilities. In this section, we show that the same is true of Sparsify. As said, this is useful because more efficient algorithms are known for estimating effective resistances than for computing them exactly. We will also use this fact for our space-bounded algorithm for sparsification in Section 5.

The lemma below says that if we only have small multiplicative approximations to the effective resistances, then the guarantees of Theorem 3.1 still hold with a small loss in the sparsity.

LEMMA 3.12. *Let $G = (V, E, w)$ be an undirected connected weighted graph on n vertices with Laplacian L . Let $0 < \varepsilon < 1$, $0 < \delta < 1/2$ and let $k \leq \log n$ be an even integer. For each $(a, b) \in E$, let \tilde{R}_{ab} be such that*

$$(1 - \alpha) \cdot R_{ab} \leq \tilde{R}_{ab} \leq (1 + \alpha) \cdot R_{ab},$$

where R_{ab} is the effective resistance of (a, b) and $0 < \alpha < 1$. Let $\tilde{R} = \{\tilde{R}_{ab}/(1 - \alpha)\}_{(a,b) \in E}$. Let H be the output of $\text{Sparsify}(G, \tilde{R}, k, \varepsilon, \delta)$ and let \tilde{L} be its Laplacian. Then, with probability at least $1 - 2\delta$, we have that

- (1) $\tilde{L} \approx_\varepsilon L$ and
- (2) H has $O\left(\frac{1+\alpha}{1-\alpha} \cdot \frac{1}{\delta^{1+2/k}} \cdot \frac{\log n}{\varepsilon^2} \cdot n^{1+\frac{2}{k}}\right)$ edges.

PROOF. Using $\tilde{R}_{ab}/(1 - \alpha)$ in Sparsify, our sampling probabilities become

$$\tilde{p}_{ab} = \min\{1, w_{ab} \cdot \tilde{R}_{ab} \cdot s/(1 - \alpha)\}.$$

This means that the expected sparsity of the resulting graph is

$$\begin{aligned} \sum_{(a,b) \in E} \tilde{p}_{ab} &\leq \sum_{(a,b) \in E} s \cdot w_{ab} \cdot \tilde{R}_{ab}/(1 - \alpha) \\ &\leq s \cdot \frac{1 + \alpha}{1 - \alpha} \cdot \sum_{(a,b) \in E} w_{ab} \cdot R_{ab} \\ &= s \cdot \frac{1 + \alpha}{1 - \alpha} \cdot (n - 1). \end{aligned}$$

Note that by feeding $\text{Sparsify } \tilde{R}_{ab}/(1 - \alpha)$ rather than \tilde{R}_{ab} , we guarantee that the approximate effective resistance is an upper bound on the true effective resistance and, hence, the approximate sampling probability is an upper bound on the true sampling probability. In particular, this implies that if $p_{ab} = 1$, then $\tilde{p}_{ab} = 1$.

Note that in Lemma 3.3 through Lemma 3.7, the expectations of X_{ab} , Z_{ab} , and \hat{Z}_{ab} do not depend on the sampling probabilities. The sampling probabilities come up when we bound two terms of the concentration bound in Theorem 3.8. However, because of our guarantee that $\tilde{p}_{ab} \geq p_{ab}$, we get the same results. The calculation we used for the first term, given in Claim 3.11, now yields

$$\begin{aligned} \mathbb{E}[\hat{Z}_{ab}^2] &= R_{ab} \cdot w_{ab}^2 \cdot \left(\frac{1}{\tilde{p}_{ab}} - 1 \right) \cdot L^{+/2} L_{ab} L^{+/2} \\ &\leq \frac{\tilde{R}_{ab}}{1 - \alpha} \cdot w_{ab}^2 \cdot \left(\frac{1}{\tilde{p}_{ab}} - 1 \right) \cdot L^{+/2} L_{ab} L^{+/2}. \end{aligned}$$

Again, when $\tilde{p}_{ab} = 1$, the above is 0; otherwise, we have that $\tilde{p}_{ab} = w_{ab} \cdot \tilde{R}_{ab} \cdot s/(1 - \alpha)$. Thus,

$$\mathbb{E}[\hat{Z}_{ab}^2] \leq w_{ab} \cdot L^{+/2} L_{ab} L^{+/2},$$

which is exactly the bound in Claim 3.11. Similarly, when adapting Claim 3.10 to the switch to \tilde{p}_{ab} , we incur no loss. We have that

$$\begin{aligned} \|\hat{Z}_{ab}\| &\leq \frac{1}{\tilde{p}_{ab}} \cdot \|w_{ab} \cdot L^{+/2} L_{ab} L^{+/2}\| \\ &= \frac{w_{ab} \cdot R_{ab}}{\tilde{p}_{ab}} \\ &\leq \frac{w_{ab} \cdot R_{ab}}{p_{ab}} = \frac{1}{s}, \end{aligned}$$

which matches the original bound.

In fact, this lemma holds with a slightly weaker assumption. Note that we used the fact that $(1 - \alpha) \cdot R_{ab} \leq \tilde{R}_{ab}$ but for the upper bound on the approximate effective resistances, we only need the weaker inequality,

$$\sum_{(a,b) \in E} w_{ab} \cdot \tilde{R}_{ab} \leq (1 + \alpha) \cdot \sum_{(a,b) \in E} w_{ab} \cdot R_{ab},$$

for the argument above to go through. \square

Note that we could equivalently define Sparsify to take approximate sampling probabilities as input (i.e., $(1 - \alpha)p_{ab} \leq \tilde{p}_{ab} \leq (1 + \alpha)p_{ab}$) rather than α -approximate effective resistances and the same lemma applies.

4 LOWER BOUNDS FOR BOUNDED-INDEPENDENCE SAMPLING

In this section, we give a lower bound for sampling-based bounded independence sparsification. Our lower bound will hold even for unweighted, simple, regular graphs in which all the effective resistances are the same. Thus, for this section, assume that $G = (V = [n], E)$ is such a graph. In Section 3, we measure sparsity in terms of the number of edges in the graph. We use this measure rather than average degree because, in weighted graphs, the degree of a vertex v typically refers to the sum of the weights of the edges incident to v , whereas in sparsification algorithms we are trying to minimize the number of edges incident to v regardless of their weight. In this section, we will sometimes refer to average degree rather than number of edges. When we refer to the average

degree of a weighted graph, we mean the average number of edges incident to each vertex. For simple, unweighted graphs, these quantities are the same.

Fix some $\alpha > 0$. Theorem 3.1 tells us that if we want to sparsify G to within error ε and expected degree $s = O(n^\alpha \cdot \log n / \varepsilon^2)$, we can do so by sampling each edge with probability $p = s \cdot (n-1)/|E|$ in a k -wise independent manner, where $k = 2/\alpha$ (rounded to an even integer).⁶ We now prove that $k \geq 4/3\alpha$ is essential for such a sampling procedure, at least for constant α .

THEOREM 4.1 (LOWER BOUND FOR SPECTRAL SPARSIFICATION VIA BOUNDED INDEPENDENCE). *Fix $c > 0$. For every $\alpha \in (0, 4/15]$, there exist infinitely many n 's for which the following holds.*

There exists a connected graph $G = (V, E)$, where $V = [n]$, whose effective resistances are all equal and a distribution $\mathcal{D} \sim \{0, 1\}^{|E|}$ that is k -wise independent for $k = \lfloor 4/3\alpha \rfloor$ with marginals $1/2$ that would fail to sparsify G to within any error $\varepsilon > 0$ and expected degree $s = c \log n \cdot n^{\alpha_0}$, where $\alpha_0 \geq (1 - 2\alpha)\alpha$.

More specifically, sampling a subgraph of G according to \mathcal{D} would result in a disconnected graph with probability at least $1 - 2/2^n$.

We note that a disconnected graph fails to be a good spectral sparsifier of a connected graph, which is implicit in Theorem 4.1. Formally:

CLAIM 4.2. *Let G and \tilde{G} be undirected graphs on n vertices with Laplacians L and \tilde{L} , respectively. If G is connected and \tilde{G} is disconnected, then $L \not\approx_\varepsilon \tilde{L}$ for any $\varepsilon > 0$.*

We give a proof of Claim 4.2 in Appendix A.

4.1 Moore-like Graphs With a Given Girth

Toward proving Theorem 4.1, we will need, for every bounded-independence parameter k , an infinite family of graphs satisfying certain properties. Recall that the *girth* of a graph G is the length of the shortest cycle in G . We will need an infinite family of girth- g graphs having large degree. Formally:

Definition 4.3. Given $\gamma > 0$ and $g: \mathbb{N} \rightarrow \mathbb{N}$, an infinite family of graphs $\{G_i = (V_i = [n_i], E_i)\}_{i \in \mathbb{N}}$ is (g, γ) -Moorish if for every $i \in \mathbb{N}$, G_i is connected, has girth at least $g(n_i)$, and is d_i -regular for some $d_i \geq n_i^{\gamma/g(n_i)} + 1$.

The problem of finding such families of graphs, or even proving their existence in some regime of parameters, has been widely studied in extremal graph theory. A simple counting argument ([20], see also [9]) shows that (g, γ) -Moorish families of graphs can only exist when $\gamma \leq 2$:

LEMMA 4.4 (THE MOORE BOUND, SEE, E.G., [9]). *Every d -regular graph of girth g on n vertices satisfies $n \geq 2 \cdot ((d-1)^{g/2} - 1)/(d-2)$.*

Still, no families with γ approaching 2 for arbitrary girths are known. The Ramanujan graphs of Lubotzky et al. [40] were shown to obtain $\gamma \geq 4/3$ by Biggs and Boshier [8]. Lazebnik et al. [34] slightly improved upon [40] in the lower-order terms. However, more importantly for us, the family they construct consists of *edge-transitive graphs*.

⁶We used the fact that for every $(a, b) \in E$, $p_{ab} \leftarrow \min \{1, R_{abs}\} = R_{abs} = R \cdot s$, which can be argued as follows. When all effective resistances equal R , we have that $R = (n-1)/|E|$ due to Claim 2.2. Now, if G has $n \cdot s$ edges or fewer, then it already achieves the desired sparsity; thus, without loss of generality, we can assume that $|E| > n \cdot s$. Hence, $R \cdot s < (n-1)s/ns < 1$. Also, the resulting graph should indeed be a weighted one; however, all its weights will be the same, $1/p$.

THEOREM 4.5 ([34]). *For every prime power d and even integer $g \geq 6$, there exists a d -regular explicit simple, edge-transitive graph with $n \leq 2d^{g-1} \lfloor \frac{g-3}{4} \rfloor - 4$ vertices and girth g . In particular, for every prime power d , there exists a $(g, \gamma = 4/3)$ -Moorish family of edge-transitive graphs, where $\text{Im}(g) = \{6, 8, \dots\}$.*

Intuitively, in an edge-transitive graph, the local environment of every edge (i.e., the vertices and edges adjacent to it) looks the same. More formally, an edge-transitive graph is one in which any two edges are equivalent under some element of its automorphism group. As the computation of the effective resistance is not affected by an automorphism, we can conclude the following claim.

CLAIM 4.6. *Let $G = (V, E)$ be an unweighted edge-transitive graph. Then, for every two edges $e = (a, b)$ and $e' = (a', b')$ in E , it holds that $R_{ab} = R_{a'b'}$.*

4.2 The Lower Bound Proof

We next prove our main result for this section, showing that Moorish edge-transitive graphs cannot be sparsified via bounded-independence edge sampling when k is too small. Our proof can be seen as an extension of an argument by Alon and Nussboim [4], who studied the bounded independence relaxation of the usual Erdős-Rényi random graph model, for which it is only required that the distribution of any subset of k edges is independent. They provide upper and lower bounds on the minimal k required to maintain properties that are satisfied by a truly random graph. In particular, they show that there exists a pairwise independent distribution \mathcal{D} over edges with marginals $1/2$ such that a random graph sampled from \mathcal{D} is disconnected almost surely.

As a warm-up, we extend the argument in [4] and show that 3-wise independence also does not suffice, even for the special case of sparsifying the complete graph.

LEMMA 4.7. *Let $G = (V = [n], E)$ be the complete graph. There exists a distribution $\mathcal{D} \sim \{0, 1\}^{|E|}$ that is 3-wise independent with marginals $1/4$ such that sampling a subgraph of G according to \mathcal{D} would result in a disconnected graph with probability at least $1 - 2/2^n$.⁷*

PROOF. We first set some notations. Let $\mathcal{G}(A, p)$ be the usual Erdős-Rényi model, in which each edge between two vertices in A is included in the graph with probability p . Let $\mathcal{B}(A)$ be the natural distribution over complete bipartite graphs: Choose a partition $A = A_1 \uplus A_2$ uniformly at random and include all edges between A_1 and A_2 .

We construct $\mathcal{D} \sim \{0, 1\}^{|E|}$ as follows. Choose a partition $[n] = V_0 \uplus V_1$ uniformly at random. On V_0 , draw a graph from $\mathcal{G}(V_0, 1/2)$ and on V_1 , draw a graph from $\mathcal{B}(V_1)$. Clearly, sampling G' according to $\mathcal{D} \sim \mathcal{D}$ would result in a disconnected graph unless $V_0 = \emptyset$ or $V_1 = \emptyset$, which occurs with probability at most $2/2^n$, so that what is left to show is that \mathcal{D} is 3-wise independent with marginals $1/4$.

This is equivalent to saying that for every $T \subseteq E$ with $|T| \leq 3$, $\Pr[\forall e \in T, D(e) = 1] = 1/4^{|T|}$. Let us first consider the case $|T| = 1$, i.e., $T = e$ for a single edge $e \in E$. Note that $\Pr[D(e) = 1]$ only if both endpoints of e appear in the same side of the partition $V_0 \uplus V_1$, which occurs with probability $1/2$ and, given that this occurs, e appears in $\mathcal{G}(V_0, 1/2)$ or $\mathcal{B}(V_1)$ with probability $1/2$. Thus, $\Pr[D(e) = 1] = 1/4$, as desired.

Next, fix a set $T \subseteq E$ of $t \in \{2, 3\}$ edges and note that we can assume without loss of generality that these edges form either a path or a triangle (for $t = 3$), as disjoint paths will occur independently. If T forms a path, then, similarly,

$$\Pr[\forall e \in T, D(e) = 1] = \Pr[V(T) \subseteq V_0] \cdot 2^{-t} + \Pr[V(T) \subseteq V_1] \cdot 2^{-t} = 2^{-(t+1)} \cdot 2^{-t} + 2^{-(t+1)} \cdot 2^{-t} = 4^{-t},$$

⁷Slightly abusing notation, by a k -wise independent distribution we mean, as customary, that the corresponding random variable is k -wise independent.

which is what we want. If T forms a triangle, then using the fact that a bipartite graph is triangle-free,

$$\Pr[\forall e \in T, D(e) = 1] = \Pr[V(T) \subseteq V_0] \cdot \frac{1}{8} = 4^{-3},$$

concluding the proof. \square

The above lemma shows that one cannot sparsify the complete graph via $(k = 3)$ -wise independent edge sampling. For a general k , we indeed need to resort to Moore-like graphs.

PROOF OF THEOREM 4.1. Recalling that $k = \lfloor 4/3\alpha \rfloor$, let $g = k + 1$ or $g = k + 2$, whichever is even. Set d_0 to be the first prime power larger than

$$1 + \max \left\{ 2^{6^2/\alpha^8}, (2c)^{6/\alpha^2} \right\}.$$

By Theorem 4.5, for every prime power $d \geq d_0$, there exists $n = n(g, d)$ and a girth- g , edge-transitive, d -regular graph $G = (V = [n], E)$ (note that, by our assumption that $\alpha \leq 4/15$, indeed $g \geq 6$). From here onwards, fix such a d and $n = n(g, d)$, observing that $\{n(g, d)\}_{d \geq d_0}$ is infinite.

Choose α_0 so that $c \log n \cdot n^{\alpha_0} = d/2$ so that marginals $1/2$ correspond to expected sparsity s . Using the fact that $d \geq n^{4/3g} + 1$, it can be verified that $\alpha_0 \geq (1 - 2\alpha) \cdot \alpha$. Indeed,

$$\alpha_0 \geq \frac{4}{3g} - \frac{\log(2c \log n)}{\log n} \geq \frac{4}{\frac{4}{\alpha} + 6} - \frac{\log(2c \log n)}{\log n} \geq \left(1 - \frac{3\alpha}{2}\right) \alpha - \frac{\log(2c \log n)}{\log n}.$$

As $n \leq (d-1)^{\frac{3g}{4}}$ and $n \geq 2 \cdot \frac{(d-1)^{\frac{g}{2}-1}}{d-2} \geq (d-1)^{\frac{g}{2}-1}$, the latter being the Moore bound, we have that

$$\begin{aligned} \frac{\log(2c \log n)}{\log n} &\leq \frac{\log(2c) + \log \frac{3g}{4} + \log \log(d-1)}{\left(\frac{g}{2} - 1\right) \log(d-1)} \\ &\leq \frac{\log(2c)}{\log(d-1)} + \frac{2}{\log(d-1)} + \frac{\log \log(d-1)}{\log(d-1)} \leq \frac{\log(2c)}{\log(d-1)} + \frac{2 \log \log(d-1)}{\log(d-1)} \leq \frac{\alpha^2}{2}, \end{aligned}$$

where we used $\log \frac{3g}{4} \leq 2(\frac{g}{2} - 1)$, $2 \leq \log \log(d-1)$, $\frac{\log(2c)}{\log(d-1)} \leq \frac{\alpha^2}{6}$, and $\frac{\log \log(d-1)}{\log(d-1)} \leq \frac{\alpha^2}{6}$. Thus, overall, $\alpha_0 \geq (1 - 2\alpha)\alpha$.

We now give a k -wise independent distribution with marginals $1/2$ that fails to yield a good spectral sparsifier for G ; namely, it will be disconnected with high probability.

To do so, construct $\mathcal{D} \sim \{0, 1\}^{|E|}$ as follows. Choose a partition $[n] = V_0 \uplus V_1$ uniformly at random. Each random partition gives rise to $D \sim \mathcal{D}$ in which for $e = (u, v) \in E$, $D(e) = 1$ (i.e., the edge e is chosen to survive) if and only if either $u, v \in V_0$ or $u, v \in V_1$.

CLAIM 4.8. *The distribution \mathcal{D} is k -wise independent with marginals $1/2$.*

PROOF. As in the proof of Lemma 4.7, it suffices to show that for every set $T \subseteq E$ of $t \leq k$ edges of G , we have that $\Pr[\forall e \in T, D(e) = 1] = 2^{-t}$. First, similar to Lemma 4.7, note that we can assume without loss of generality that T is a connected component since whenever T_1 and T_2 are over disjoint sets of vertices, $\Pr[\forall e \in T_1 \cup T_2, D(e) = 1] = \Pr[\forall e \in T_1, D(e) = 1] \cdot \Pr[\forall e \in T_2, D(e) = 1]$. As the girth of G is larger than t , it must be the case that A is a tree.

In such a case, where T contains no cycles, $\Pr[\forall e \in T, D(e) = 1]$ is equal to the probability that all $t + 1$ vertices in T belong to the same partition, which is $2 \cdot 2^{-(t+1)} = 2^{-t}$. \square

By the way that \mathcal{D} was constructed, it is clear that sampling G' according to \mathcal{D} would result in a disconnected graph unless $V_0 = \emptyset$ or $V_1 = \emptyset$, which occurs with probability $1 - 2/2^n$, meaning that G' almost surely does not ε -approximate G for any ε . \square

We again stress that by the work in Section 3, we know that *any* k -wise independent distribution over the edges of G with marginals $s \cdot (n-1)/|E|$ for $k = \lceil 2/\alpha \rceil$ would produce an ε -spectral sparsifier with expected degree $O(s)$ with high constant probability.

The above also implies that any improvement upon Moorish families of edge-transitive graphs will improve our lower bound. Assuming the existence of a $(g, \gamma = 2)$ -Moorish family of edge-transitive graphs, we are able to show that the result of Section 3 is essentially tight.

5 SPECTRAL SPARSIFIERS IN DETERMINISTIC SMALL SPACE

In this section, we show that Sparsify can be derandomized space efficiently.

THEOREM 5.1 (DETERMINISTIC SMALL-SPACE SPARSIFICATION). *Let G be an undirected, connected, weighted graph on n vertices with Laplacian L . There is a deterministic algorithm that, when given G , an even integer k and $0 < \varepsilon < 1$ outputs a weighted graph H with Laplacian \tilde{L} satisfying:*

- (1) $\tilde{L} \approx_{\varepsilon} L$ and
- (2) H has $O(\frac{\log n}{\varepsilon^2} n^{1+2/k})$ edges.

The algorithm runs in space $O(k \log(N \cdot w) + \log(N \cdot w) \log \log(N \cdot w))$, where $w = w_{\max}/w_{\min}$ is the ratio of the maximum and minimum edge weights in G and N is the bitlength of the input.

We use the standard model of space-bounded computation. The machine has a read-only input tape, a constant number of read/write work tapes, and a write-only output tape. We say the machine runs in space s if, throughout the computation, it only uses s total tape cells on the work tapes. The machine may write outputs to the output tape that are larger than s (in fact, as large as $2^{O(s)}$) but the output tape is write-only. We use the following fact about the composition of space-bounded algorithms (see [5, Chapter 4] or [23, Chapter 5]).

LEMMA 5.2. *Let f_1 and f_2 be functions that can be computed in space $s_1(n), s_2(n) \geq \log n$, respectively, and f_1 has output of length $\ell_1(n)$ on inputs of size n . Then, $f_2 \circ f_1$ can be computed in space*

$$O(s_2(\ell_1(n)) + s_1(n)).$$

The natural way to derandomize Sparsify would be to iterate over all elements of the corresponding k -wise independent sample space. More formally, given $\{p_{ab}\}_{(a,b) \in E}$, let I_{ab} be the indicator random variable that is 1 if and only if edge (a, b) is chosen. If the I_{ab} s are k -wise independent so that $\Pr[I_{ab} = 1] = p_{ab}$ (or some good approximation of p_{ab}), we are guaranteed to succeed with non-zero probability. Hence, at least one assignment to the I_{ab} s taken from the k -wise independent is guaranteed to work. From Section 2.4, we know that the sample space is small enough that we can afford to enumerate over all elements in it. Towards proving Theorem 5.1, there are still three issues to consider:

- (1) Approximating the effective resistances R_{ab} for every $(a, b) \in E$, space efficiently. Fortunately, we can do this with high accuracy using the result of Murtaghet al. [43] for approximating the pseudoinverse of a Laplacian, which we state shortly.
- (2) Verifying that a given set of random choices in Sparsify provides a sparse and accurate approximation to the input graph. The sparsity requirement is easy to check. To check that $\tilde{L} \approx_{\varepsilon} L$, we devise a verification algorithm that uses the algorithm of [43]. The details are given in Lemma 5.7.
- (3) The Laplacian solver of [43] only works for multigraphs (graphs with integer edge weights) and we want an algorithm that works for general weighted graphs. To fix this, we extend the work of [43] by giving a simple reduction from the weighted case to the multigraph case. The details can be found in Appendix B.

5.1 Algorithm for Approximating Effective Resistances

A key ingredient in our deterministic sparsification algorithm is a deterministic nearly logarithmic space algorithm for approximating the pseudoinverse of an undirected Laplacian.

THEOREM 5.3 ([43]). *Given an undirected, connected multigraph G with Laplacian $L = D - A$ and $\varepsilon > 0$, there is a deterministic algorithm that computes a symmetric PSD matrix \widetilde{L}^+ such that $\widetilde{L}^+ \approx_\varepsilon L^+$ and uses space $O(\log N \cdot \log \log \frac{N}{\varepsilon})$, where N is the bitlength of the input (as a list of edges).*

Note that the space complexity above assumes that the multigraph is given as a list of edges. If we instead think of parallel edges as integer edge weights, then N should be replaced by $N \cdot w_{\max}$ where w_{\max} is the maximum edge weight in G since an edge of weight w gets repeated w times in the edge-list representation. To work with general weighted graphs, we extend the result of [43].

LEMMA 5.4 (SMALL SPACE LAPLACIAN SOLVER FOR WEIGHTED GRAPHS). *Given an undirected connected weighted graph $G = (V, E, w)$ with Laplacian $L = D - A$, and $0 < \varepsilon < 1$, there exists a deterministic algorithm that computes a symmetric PSD matrix \widetilde{L}^+ such that $\widetilde{L}^+ \approx_\varepsilon L^+$ and uses space $O(\log(N \cdot w) \log \log(N \cdot w/\varepsilon))$, where $w = w_{\max}/w_{\min}$ is the ratio of the maximum and minimum edge weights in G and N is the bitlength of the input.*

A proof of Lemma 5.4 can be found in Appendix B. Lemma 5.4 immediately gives an algorithm for computing strong multiplicative approximations to effective resistances.

LEMMA 5.5. *Let $G = (V, E, w)$ be an undirected, connected, weighted graph and let R_{ab} be the effective resistance of $(a, b) \in E$. There is an algorithm that computes a real number \widetilde{R}_{ab} such that*

$$(1 - \varepsilon) \cdot R_{ab} \leq \widetilde{R}_{ab} \leq (1 + \varepsilon) \cdot R_{ab}$$

and uses space $O(\log(N \cdot w) \cdot \log \log \frac{N \cdot w}{\varepsilon})$, where $w = w_{\max}/w_{\min}$ is the ratio of the maximum and minimum edge weights in G and N is the bitlength of the input.

PROOF. Let L be the Laplacian of G . By the definition of effective resistance, we have that $R_{ab} = (e_a - e_b)^\top L^+ (e_a - e_b)$. From Lemma 5.4, we can compute a matrix \widetilde{L} such that

$$(1 - \varepsilon) \cdot L^+ \leq \widetilde{L}^+ \leq (1 + \varepsilon) \cdot L^+$$

in space $O(\log(N \cdot w) \cdot \log \log \frac{N \cdot w}{\varepsilon})$. By the definition \leq , this implies that

$$(1 - \varepsilon) \cdot (e_a - e_b)^\top L^+ (e_a - e_b) \leq (e_a - e_b)^\top \widetilde{L}^+ (e_a - e_b) \leq (1 + \varepsilon) \cdot (e_a - e_b)^\top L^+ (e_a - e_b).$$

Setting $\widetilde{R}_{ab} = (e_a - e_b)^\top \widetilde{L}^+ (e_a - e_b)$ and noting that the vector matrix multiplication only adds logarithmic space overhead completes the proof. \square

5.2 Testing for Spectral Proximity

In this section, we give our deterministic, small-space procedure for verifying that two Laplacians spectrally approximate one another. We will need the following claim about the space complexity of matrix multiplication.

CLAIM 5.6. *Given $n \times n$ matrices M_1, \dots, M_k , their product $M_1 \cdot \dots \cdot M_k$ can be computed using $O(\log N \cdot \log k)$ space, where N is the bitlength of (M_1, \dots, M_k) .*

The proof of Claim 5.6 uses the natural divide-and-conquer algorithm and the fact that two matrices can be multiplied in logarithmic space. A detailed proof can be found in [43].

Using Lemma 5.4 and Claim 5.6, we prove the following lemma. The high-level idea is that testing whether two matrices L and \tilde{L} spectrally approximate each other can be reduced to approximating the spectral radius of a particular matrix

$$M = \left(\frac{(\tilde{L} - L)L^+}{\varepsilon} \right)^2.$$

In fact, it will be sufficient to check whether the trace of a sufficiently high power of M is below a certain threshold to deduce whether the spectral radius of M does not exceed 1. For intuition, replace the matrices with scalars m , ℓ , and $\tilde{\ell}$, where

$$m = \frac{(\tilde{\ell} - \ell)^2}{(\varepsilon \cdot \ell)^2}.$$

Then, $m \leq 1$ implies that $\sqrt{m} \leq 1$, which implies that $|\tilde{\ell} - \ell| \leq \varepsilon \cdot \ell$ — the kind of relative closeness we want between the matrices \tilde{L} and L when aiming for spectral approximation.

LEMMA 5.7. *There exists a deterministic algorithm that, given undirected, connected, weighted graphs \tilde{G} and G with Laplacians \tilde{L} , L , and $\varepsilon, \alpha > 0$, outputs YES or NO such that*

- (1) *If $\tilde{L} \approx_\varepsilon L$, then the algorithm outputs YES and*
- (2) *If $\tilde{L} \not\approx_{\varepsilon \cdot \sqrt{1+\alpha}} L$, then the algorithm outputs NO.*

The algorithm uses space $O(\log(N \cdot w) \cdot \log \log \frac{N \cdot w}{\alpha \varepsilon} + \log(N \cdot w) \cdot \log \frac{1}{\alpha})$, where $w = w_{\max}/w_{\min}$ is the ratio of the maximum and minimum edge weights in G and \tilde{G} and N is the bitlength of the input.

PROOF. Let

$$M = \left(\frac{(\tilde{L} - L)L^+}{\varepsilon} \right)^2.$$

Set $T = \text{Tr}(M)$ and $t = \lceil \frac{\log T}{\log(1+\alpha)} \rceil$. The following claim shows that if we can compute $\text{Tr}(M^t)$ exactly, then we can check the two cases in Lemma 5.7. However, we won't be able to compute $\text{Tr}(M^t)$ exactly because that would require computing L^+ exactly. This will be addressed later.

CLAIM 5.8. *If $\tilde{L} \approx_\varepsilon L$, then $\text{Tr}(M^t) \leq T$ and if $\tilde{L} \not\approx_{\varepsilon \cdot \sqrt{1+\alpha}} L$, then $\text{Tr}(M^t) > T$.*

PROOF. Let $\Pi = I - J$ be the orthogonal projection onto $\text{Im}(L) = \text{Im}(\tilde{L})$ (note that both L and \tilde{L} are the Laplacians of connected graphs). Using Claims 2.7 and 2.8, we know that $\tilde{L} \approx_\varepsilon L$ if and only if for all $v \in \mathbb{R}^n$ we have that

$$-v^\top \Pi v \leq \frac{1}{\varepsilon} \cdot v^\top L^{+/2} (\tilde{L} - L) L^{+/2} v \leq v^\top \Pi v.$$

Note that Π and $L^{+/2}(\tilde{L} - L)L^{+/2}$ have the same kernel — namely, $\text{span}(\{1\})$ — and being perpendicular to 1 is preserved under both operators. Thus, the above holds if and only if it holds on all vectors $v \perp 1$. For such vectors, we have that $\Pi v = v$; hence, $v^\top \Pi v = \|v\|^2$. Thus, we have that $\tilde{L} \approx_\varepsilon L$ if and only if for all vectors $v \perp 1$ we have that

$$\left| \frac{1}{\varepsilon} \cdot v^\top L^{+/2} (\tilde{L} - L) L^{+/2} v \right| \leq \|v\|^2$$

or, equivalently,

$$\rho \left(\frac{1}{\varepsilon} \cdot L^{+/2} (\tilde{L} - L) L^{+/2} \right) \leq 1.$$

Note that $L^{+1/2}(\tilde{L} - L)L^{+1/2}$ is symmetric; thus, it has real eigenvalues and is similar to the matrix $(\tilde{L} - L)L^+$ on the space orthogonal to the kernel. Therefore, we can rewrite the above condition as

$$\rho\left(\frac{1}{\varepsilon} \cdot (\tilde{L} - L)L^+\right) \leq 1.$$

Furthermore, we have that $M = ((\tilde{L} - L)L^+/\varepsilon)^2$ has real, non-negative eigenvalues.

Consider the contrapositives of the implications stated in the claim, namely: If $\text{Tr}(M^t) > T$, then $\tilde{L} \not\approx_\varepsilon L$ and if $\text{Tr}(M^t) \leq T$, then $\tilde{L} \approx_{\varepsilon \cdot \sqrt{1+\alpha}} L$. Now, note that if $\text{Tr}(M^t) > T = \text{Tr}(M)$, then $\rho(M) > 1$ because the only way that the trace of a matrix with real non-negative eigenvalues can increase under powering is if at least one of its eigenvalues exceeds 1. Thus, on the one hand, if $\text{Tr}(M^t) > T = \text{Tr}(M)$, then $\tilde{L} \not\approx_\varepsilon L$. On the other hand,

$$\begin{aligned} \text{Tr}(M^t) \leq T &\implies \rho(M)^t \leq T \\ &\implies \rho(M) \leq T^{1/t} \leq (1 + \alpha). \end{aligned}$$

When $\rho(M) \leq 1 + \alpha$, we have that $\rho(L^{+1/2}(\tilde{L} - L)L^{+1/2}/\varepsilon) \leq \sqrt{1 + \alpha}$; thus,

$$-\sqrt{1 + \alpha} \cdot \Pi \leq \frac{1}{\varepsilon} \cdot L^{+1/2}(\tilde{L} - L)L^{+1/2} \leq \sqrt{1 + \alpha} \cdot \Pi$$

or, in other words, $\tilde{L} \approx_{\varepsilon \cdot \sqrt{1+\alpha}} L$. □

Recall that we cannot compute M exactly in small space because we do not know how to compute L^+ exactly in small space. Let

$$\widehat{M} = \left(\frac{1}{1 + \gamma} \cdot \frac{(\tilde{L} - L)\widehat{L}}{\varepsilon} \right)^2,$$

where $\widehat{L} \approx_\gamma L^+$ for

$$\gamma = 1 - \frac{2}{1 + \sqrt{1 + \frac{\alpha}{2 + \alpha}}}.$$

We have chosen γ so that $\left(\frac{1+\gamma}{1-\gamma}\right)^2 = 1 + \frac{\alpha}{2+\alpha}$; hence,

$$\left(\frac{1+\gamma}{1-\gamma}\right)^2 \cdot \left(1 + \frac{\alpha}{2}\right) = 1 + \alpha.$$

Now, we will show that \widehat{M} is sufficient for our purposes. For this, we will need the following claim, whose proof is deferred to Appendix A.

CLAIM 5.9. *Let C be a real, symmetric matrix and let A and B be real, symmetric, PSD matrices such that $A \leq B$. Suppose that $\ker(A) = \ker(B) = \ker(C)$. Then,*

$$\rho(CA) \leq \rho(CB).$$

Next, we show that \widehat{M} is sufficient for our purposes.

CLAIM 5.10. *It holds that $\left(\frac{1-\gamma}{1+\gamma}\right)^2 \cdot \rho(M) \leq \rho(\widehat{M}) \leq \rho(M)$.*

This claim is sufficient because we can use the procedure implied by Claim 5.8 to distinguish the case of $\rho(\widehat{M}) \leq 1$ from the case of $\rho(\widehat{M}) > 1 + \alpha'$, for $\alpha' = \frac{\alpha}{2}$. From Claim 5.10, we know that

$$\rho(M) \leq 1 \implies \rho(\widehat{M}) \leq 1$$

and that

$$\rho(M) > 1 + \alpha \implies \rho(\widehat{M}) > \left(\frac{1 - \gamma}{1 + \gamma} \right)^2 \cdot (1 + \alpha) = 1 + \alpha'.$$

Thus, from the arguments above, we can distinguish the case of $\widetilde{L} \approx_\varepsilon L$ from $\widetilde{L} \not\approx_{\varepsilon \cdot \sqrt{1+\alpha}} L$ by computing $\text{Tr}(\widehat{M}^{\widehat{t}})$ for $\widehat{t} = \lceil \log \text{Tr}(\widehat{M}) / \log(1 + \alpha') \rceil$ and comparing the result to $\text{Tr}(\widehat{M})$. Now, we prove Claim 5.10.

PROOF. By assumption, we have that $\widehat{L} \approx_\gamma L^+$; hence,

$$\frac{1 - \gamma}{1 + \gamma} \cdot L^+ \leq \frac{1}{1 + \gamma} \cdot \widehat{L} \leq L^+.$$

Let $C = \frac{1}{\varepsilon}(\widetilde{L} - L)$ and note that C is symmetric. By Claim 5.9, we have that

$$\frac{1 - \gamma}{1 + \gamma} \cdot \rho(CL^+) \leq \frac{1}{1 + \gamma} \cdot \rho(\widehat{CL}) \leq \rho(CL^+)$$

and, hence, that

$$\left(\frac{1 - \gamma}{1 + \gamma} \cdot \rho(CL^+) \right)^2 \leq \left(\frac{1}{1 + \gamma} \cdot \rho(\widehat{CL}) \right)^2 \leq \rho(CL^+)^2.$$

Noting that for all matrices A with real eigenvalues, we have that $\rho(A^2) = \rho(A)^2$ and that $(CL^+)^2 = M$ and $(\widehat{CL}/(1 + \gamma))^2 = \widehat{M}$, the above becomes

$$\left(\frac{1 - \gamma}{1 + \gamma} \right)^2 \cdot \rho(M) \leq \rho(\widehat{M}) \leq \rho(M),$$

as desired. \square

Thus, our distinguishing algorithm goes as follows: Approximate $\widehat{L} \approx_\gamma L^+$ using the Laplacian solver algorithm given in Lemma 5.4, compute $\text{Tr}(\widehat{M}^{\widehat{t}})$ and answer according to whether it is greater than or less than $\text{Tr}(\widehat{M})$. Claims 5.8 and 5.10 establish its correctness. We are left with establishing the space complexity.

CLAIM 5.11. *The distinguishing algorithm uses space $O(\log(N \cdot w) \cdot \log \log \frac{N \cdot w}{\alpha \varepsilon} + \log(N \cdot w) \cdot \log \frac{1}{\alpha})$, where $w = w_{\max}/w_{\min}$ is the ratio of the maximum and minimum edge weights in G and \widetilde{G} and N is the bitlength of the input.*

PROOF. By Lemma 5.4, we can compute \widehat{L} in space $S = O(\log(N \cdot w) \log \log \frac{N \cdot w}{\gamma}) = O(\log(N \cdot w) \cdot \log \log \frac{N \cdot w}{\alpha})$, and up to constant factors, this is also the space required to compute \widehat{M} . Note that the bitlength of \widehat{M} is $N' = O(N + \log(1/\varepsilon)) = O(N)$. Claim 5.6 and composition of space-bounded algorithms (Proposition 5.2)⁸ say that we can compute $\widehat{M}^{\widehat{t}}$ using

$$S' = O\left(S + \log N' \cdot \log \widehat{t}\right) = O\left(S + \log N' \cdot \left(\log \log \text{Tr}(\widehat{M}) + \log \frac{1}{\log(1 + \alpha)}\right)\right)$$

space. Finally, computing $\text{Tr}(\widehat{M}^{\widehat{t}})$ requires summing n entries in $\widehat{M}^{\widehat{t}}$; thus, the required additional space can be bounded by $O(\log n + \log(N' \cdot \widehat{t})) = O(S')$.

Now, we argue that we can (loosely) bound $\text{Tr}(\widehat{M})$ by $\text{poly}(N \cdot w, 1/\alpha, 1/\varepsilon)$. If d_{\max} is the maximum weighted degree of the graph corresponding to L and \widetilde{L} , then we have that the spectral norms of L and \widetilde{L} are at most $2 \cdot d_{\max} = O(N \cdot w_{\max})$. Claim 2.9 says that the smallest

⁸Using the notation of Proposition 5.2, f_1 computes \widehat{M} and f_2 computes matrix powering.

non-zero eigenvalue of L is lower bounded by w_{\min}/n^2 . Note that $\|L^+\|$ equals the reciprocal of the smallest non-zero eigenvalue of L ; hence, we have that $\|L^+\| \leq n^2/w_{\min}$ and, therefore, $\|\widehat{L}\| \leq (1+\gamma) \cdot n^2/w_{\min} = \text{poly}(N/w_{\min}, 1/\alpha)$. It follows that $\text{Tr}(\widehat{M}) = \text{poly}(N \cdot w, 1/\alpha, 1/\varepsilon)$. Plugging this into the space complexity gives us a space bound of

$$O\left(\log(N \cdot w) \cdot \left(\log \log \frac{N \cdot w}{\alpha \varepsilon} + \log \frac{1}{\log(1+\alpha)}\right)\right).$$

Note that $\log \frac{1}{\log(1+\alpha)} = O(\log \frac{1}{\alpha})$ completes the analysis. \square

\square

5.3 Completing the Proof of Theorem 5.1

We can now prove Theorem 5.1. As noted above, the algorithm proceeds by first approximating the sampling probabilities and then sparsifying G where the surviving edges are chosen from a small k -wise independent sample space whose marginals are set properly. Each potential sparsifier is checked using the algorithm given in Section 5.2.

PROOF OF THEOREM 5.1. Set $\delta = \frac{1}{4}$, $\widehat{\varepsilon} = \frac{4\varepsilon}{5}$ and

$$s = \frac{18e \log n}{\widehat{\varepsilon}^2} \cdot \left(\frac{n}{\delta}\right)^{2/k}$$

for α soon to be determined. These parameters are chosen in accordance with the parameters required for Sparsify to succeed with probability $1/2$ and approximation error $\widehat{\varepsilon}$ (see Lemma 3.12). Set $\alpha' = \alpha/(4 + \alpha)$. We compute approximate effective resistances \widetilde{R}_{ab} for each edge (a, b) in G using Lemma 5.5, so that

$$(1 - \alpha')R_{ab} \leq \widetilde{R}_{ab} \leq (1 + \alpha')R_{ab}.$$

This takes $O(\log(N \cdot w) \log \log((N \cdot w)/\alpha))$ space. Then, we compute approximate sampling probabilities as follows:

$$\widetilde{p}_{ab} = \alpha' \cdot \left\lfloor \frac{1}{\alpha'} \cdot \min \left\{ 1, w_{ab} \cdot \widetilde{R}_{ab} \cdot s / (1 - \alpha') \right\} \right\rfloor$$

That is, we truncate the required (approximate) sampling probabilities to $\log \frac{1}{\alpha'}$ bits of precision. In particular, denoting the precise sampling probabilities by $p_{ab}^* = \min \{1, w_{ab} \cdot R_{ab} \cdot s\}$, we have that

$$\begin{aligned} \min \{1, w_{ab} \cdot \widetilde{R}_{ab} \cdot s / (1 - \alpha')\} - p_{ab}^* &\leq w_{ab} \cdot s \cdot R_{ab} \cdot \left(\frac{1 + \alpha'}{1 - \alpha'} - 1 \right) \\ &= p_{ab}^* \cdot \frac{2\alpha'}{1 - \alpha'} \\ &\leq \alpha/2. \end{aligned}$$

Furthermore, we have an additional error of α' due to the truncation; thus, $|\widetilde{p}_{ab} - p_{ab}^*| \leq \alpha/2 + \alpha' \leq \alpha$.

We want to set α so that \widetilde{p}_{ab} is a multiplicative approximation to p_{ab}^* for all $(a, b) \in E$, which requires α to be smaller than $\min_{(a, b) \in E} \{p_{ab}^*\}$.

CLAIM 5.12. *Let d_{\max} be the maximum weighted degree over all vertices in G . Then, for all $(a, b) \in E$, $p_{ab}^* \geq 1/d_{\max}$.*

PROOF. Since $s > 1$ and $w_{ab} \geq 1$ (all edge weights are positive integers) we have that $p_{ab}^* \geq R_{ab}$. Let $\lambda_{\min}(C)$ denote the minimal non-zero eigenvalue of a matrix C . To lower bound R_{ab} , we use the variational characterization of eigenvalues and the definition of effective resistance to write

$$\begin{aligned} R_{ab} &= (e_a - e_b)^\top L^+ (e_a - e_b) \\ &\geq \lambda_{\min}(L^+) \cdot \|e_a - e_b\|^2 \\ &= \frac{2}{\|L\|} \\ &\geq \frac{1}{d_{\max}}. \end{aligned}$$

Note that we can indeed consider the minimal non-zero eigenvalue of L^+ because $e_a - e_b$ is perpendicular to the one-dimensional kernel of L (the all-ones vector). \square

In light of the above, we can set α so that $1/\alpha = 2 \cdot d_{\max} = O(N \cdot w)$ and get a $1/2$ -multiplicative approximation to the sampling probabilities.

Now, consider the k -wise independent sample space $\mathcal{D} \subseteq \{0, 1\}^{|E|}$ guaranteed to us by Lemma 2.4, substituting $t = \lceil \log(1/\alpha') \rceil$. By Lemma 2.4, each element of \mathcal{D} can be sampled using

$$O(k \cdot \max\{\log(1/\alpha'), \log |E|\}) = O(k \cdot \log(N \cdot w))$$

space. For each element of \mathcal{D} , construct the corresponding sparse graph. Note that the space used to cycle through each element can be reused. Lemma 3.12 tells us that at least $1 - 2\delta = 1/2$ of the Laplacians of the resulting graphs $\widehat{\varepsilon}$ -approximate the Laplacian of G and have

$$O\left(\frac{1+1/2}{1-1/2} \cdot \frac{1}{\delta^{1+2/k}} \cdot \frac{\log n}{\varepsilon^2} \cdot n^{1+\frac{2}{k}}\right) = O\left(\frac{\log n}{\varepsilon^2} \cdot n^{1+\frac{2}{k}}\right)$$

edges. For each of these graphs, we run the verification algorithm with accuracy parameter $9/16$, which is guaranteed to find a graph with the above sparsity whose Laplacian approximates the Laplacian of G with error

$$\widehat{\varepsilon} \cdot \sqrt{1 + \frac{9}{16}} = \frac{4\varepsilon}{5} \cdot \sqrt{\frac{25}{16}} = \varepsilon$$

in space

$$O\left(\log(N \cdot w) \log \log \frac{16N \cdot w}{9\widehat{\varepsilon}} + \log(N \cdot w) \log \frac{16}{9}\right) = O\left(\log(N \cdot w) \log \log \frac{N \cdot w}{\varepsilon}\right).$$

Again, the space used for the verification process can be reused. Adding up the space complexities gives us a total of

$$O\left(k \log(N \cdot w) + \log(N \cdot w) \log \log \frac{N \cdot w}{\varepsilon}\right)$$

space. Note that the final result is vacuous when $\varepsilon \leq 1/n$. Thus, we can without loss of generality assume that $\varepsilon \geq 1/n$. This gives a total space complexity of $O(k \log(N \cdot w) + \log(N \cdot w) \log \log(N \cdot w))$. \square

APPENDICES

A DEFERRED PROOFS

Claim 2.5 restated. Let A, B, C be $n \times n$ symmetric PSD matrices and suppose that $B \leq C$. Then,

$$\|A + B\| \leq \|A + C\|.$$

PROOF. For every symmetric matrix M , we have that $\|M\| = \max_{x: \|x\|=1} x^\top M x$. Thus,

$$\|A + B\| = \max_{x: \|x\|=1} (x^\top A x + x^\top B x) \leq \max_{x: \|x\|=1} (x^\top A x + x^\top C x) = \|A + C\|.$$

□

Claim 2.6 restated. Let $G = (V, E, w)$ be an undirected weighted graph on n vertices with Laplacian L . Fix $(a, b) \in E$ and recall that $L_{ab} = (e_a - e_b)(e_a - e_b)^\top$. Then,

$$\|L^{+1/2} L_{ab} L^{+1/2}\| = R_{ab}.$$

PROOF. L_{ab} is PSD and has rank 1. Thus, it follows that $L^{+1/2} L_{ab} L^{+1/2}$ is PSD and has rank 1. Since the trace is the sum of the eigenvalues and the norm is the maximum eigenvalue, this implies that $\|L^{+1/2} L_{ab} L^{+1/2}\| = \text{Tr}(L^{+1/2} L_{ab} L^{+1/2})$. Finally, the trace is invariant under cyclic permutations; thus, we have that

$$\begin{aligned} \text{Tr}\left(L^{+1/2}(e_a - e_b)(e_a - e_b)^\top L^{+1/2}\right) &= \text{Tr}\left((e_a - e_b)^\top L^+(e_a - e_b)\right) \\ &= R_{ab}. \end{aligned}$$

□

Claim 2.7 restated. Let $G = (V, E, w)$ be an undirected connected weighted graph on n vertices with Laplacian L . Let J be the $n \times n$ matrix with $1/n$ in every entry and define $\Pi = I - J$. Then, we have that

$$\Pi = LL^+ = L^+L = L^{+1/2}LL^{+1/2},$$

where Π is the orthogonal projection onto the image of L (the space orthogonal to the all-ones vector).

PROOF. It is clear that for all $v = \alpha \mathbf{1}$ where $\alpha \in \mathbb{R}$, we have that $\Pi v = v - Jv = \mathbf{0}$ and for all $v \perp \mathbf{1}$, we have that $\Pi v = v$. Thus, Π is the orthogonal projection onto the space orthogonal to the all-ones vector. Since G is connected, $\ker(L) = \text{span}(\{\mathbf{1}\})$ and for every symmetric matrix M , $\text{Im}(M) = \ker(M)^\perp$, it follows that Π is also the orthogonal projection onto the image of L .

By the definition of the pseudoinverse, we have that L and L^+ have identical eigenvectors. Letting $v_1 = \mathbf{1}, v_2, \dots, v_n$ be their orthonormal basis of eigenvectors, $\lambda_1 = 0, \dots, \lambda_n$ be the eigenvalues of L and $0, \frac{1}{\lambda_2}, \dots, \frac{1}{\lambda_n}$ be the eigenvalues of L^+ , we have that

$$LL^+ v_1 = \mathbf{0}$$

and for each $i \in \{2, \dots, n\}$,

$$LL^+ v_i = \frac{1}{\lambda_i} L v_i = \frac{\lambda_i}{\lambda_i} v_i = v_i.$$

Thus, LL^+ and Π are identical on a set of basis vectors and, hence, are the same matrix. Similar calculations yield the same result for L^+L and $L^{+1/2}LL^{+1/2}$. □

Claim 2.8 restated. Let A, B, C be symmetric $n \times n$ matrices and suppose that A and B are PSD. Then, the following hold:

- (1) $A \approx_\varepsilon B \implies C^\top A C \approx_\varepsilon C^\top B C$
- (2) If $\ker(C) \subseteq \ker(A) = \ker(B)$, then $A \approx_\varepsilon B \iff C^\top A C \approx_\varepsilon C^\top B C$

PROOF. First, we prove item (1). Note that since $A \approx_\varepsilon B$, we have that $\ker(A) = \ker(B)$; otherwise, the spectral inequalities can be violated by a vector in the kernel of one of the matrices but not in the kernel of the other. Fix $x \in \mathbb{R}^n$ and let $y = Cx$. By assumption, we have that

$$(1 - \varepsilon) \cdot y^\top B y \leq y^\top A y \leq (1 + \varepsilon) \cdot y^\top B y.$$

Observing that $y^\top Ay = x^\top C^\top ACx$ and $y^\top By = x^\top C^\top BCx$ and noting that x was arbitrary completes the proof.

For item 2, assume that $\ker(C) \subseteq \ker(A) = \ker(B)$ and we will show that

$$C^\top AC \approx_\varepsilon C^\top BC \implies A \approx_\varepsilon B.$$

The other direction follows from item 1. Fix $x \in \mathbb{R}^n$. We want to show that

$$(1 - \varepsilon) \cdot x^\top Bx \leq x^\top Ax \leq (1 + \varepsilon) \cdot x^\top Bx. \quad (2)$$

If $x \in \ker(A) = \ker(B)$, then the above is trivially true. Thus, without loss of generality, we can take $x \in \ker(A)^\perp = \text{Im}(A) = \text{Im}(B)$. Let $y = C^\dagger x$. By assumption, we have that

$$(1 - \varepsilon) \cdot y^\top C^\top BCy \leq y^\top C^\top ACy \leq (1 + \varepsilon) \cdot y^\top C^\top BCy$$

Let $\Pi = CC^\dagger$. We can rewrite the above as

$$(1 - \varepsilon) \cdot x^\top \Pi^\top B \Pi x \leq x^\top \Pi^\top A \Pi x \leq (1 + \varepsilon) \cdot x^\top \Pi^\top B \Pi x.$$

Note that by the definition of the Moore–Penrose pseudoinverse, Π is the projection onto $\ker(C)^\perp \supseteq \text{Im}(A) = \text{Im}(B)$. Since we assumed without loss of generality that $x \in \text{Im}(A) = \text{Im}(B)$, it follows that $\Pi x = x$. Substituting this into the above establishes Equation (2) and completes the proof. \square

Claim 4.2 restated. *Let G and \tilde{G} be undirected graphs on n vertices with Laplacians L and \tilde{L} , respectively. If G is connected and \tilde{G} is disconnected, then $L \not\approx_\varepsilon \tilde{L}$ for any $\varepsilon > 0$.*

PROOF. Since G is a connected, undirected graph, we have that $\ker(L) = \text{span}(\{1\})$. We will show that there is a vector $v \in \ker(\tilde{L})$ such that $v \notin \ker(L)$. This will complete the proof because it implies that $v^\top \tilde{L}v = 0$ but $v^\top Lv \neq 0$ and, hence, the quadratic forms of these Laplacians cannot multiplicatively approximate each other.

Let C be a connected component of \tilde{G} and let $V \setminus C$ be the remaining vertices. By assumption, $V \setminus C \neq \emptyset$. Let $\tilde{L}(C)$ be the Laplacian of \tilde{G} if all of the edges in $V \setminus C$ were deleted and define $\tilde{L}(V \setminus C)$ analogously. Note that $\tilde{L} = \tilde{L}(C) + \tilde{L}(V \setminus C)$.

Define v so that $v_i = 0$ for all $i \in V \setminus C$ and $v_i = 1$ for all $i \in C$. Then, we have that

$$v^\top \tilde{L}v = v^\top \tilde{L}(C)v + v^\top \tilde{L}(V \setminus C)v = 0,$$

but $v \notin \ker(L)$. \square

Claim 5.9 restated. *Let C be a real, symmetric matrix and let A and B be real, symmetric, PSD matrices such that $A \leq B$. Suppose that $\ker(A) = \ker(B) = \ker(C)$. Then,*

$$\rho(CA) \leq \rho(CB).$$

PROOF. We will show that $\rho(A^{1/2}CA^{1/2}) \leq \rho(B^{1/2}CB^{1/2})$. This will complete the proof because $A^{1/2}CA^{1/2}$ is similar to CA on the space orthogonal to $\ker(A) = \ker(C)$ (multiply on the left by $A^{+1/2}$ and on the right by $A^{1/2}$). Therefore, $A^{1/2}CA^{1/2}$ and CA have the same spectrum and likewise that $B^{1/2}CB^{1/2}$ and CB have the same spectrum. Thus, establishing the inequality would show that $\rho(CA) \leq \rho(CB)$, as desired. To show that $\rho(A^{1/2}CA^{1/2}) \leq \rho(B^{1/2}CB^{1/2})$, we prove the following claim.

CLAIM A.1. *Let $M \in \mathbb{R}^{n \times n}$ be a (possibly asymmetric) matrix and let $A, B \in \mathbb{R}^{n \times n}$ be symmetric PSD matrices such that $A \leq B$. Then,*

- (1) $\|A^{1/2}M\| \leq \|B^{1/2}M\|$,
- (2) $\|MA^{1/2}\| \leq \|MB^{1/2}\|$.

PROOF. To prove item (1), we write

$$\begin{aligned}
 \|A^{1/2}M\| &= \max_{x: \|x\|=1} \|A^{1/2}Mx\| \\
 &= \max_{x: \|x\|=1} \sqrt{x^\top M^\top A M x} \\
 &\leq \max_{x: \|x\|=1} \sqrt{x^\top M^\top B M x} \\
 &= \|B^{1/2}M\|,
 \end{aligned}$$

where the inequality follows from the assumption that $A \leq B$. For the second item, first use item (1) to get $\|A^{1/2}M^\top\| \leq \|B^{1/2}M^\top\|$. Recall that for all matrices N , we have that $\|N\| = \|N^\top\|$. Therefore, since A and B are symmetric, we get that $\|MA^{1/2}\| \leq \|MB^{1/2}\|$. \square

Now, set $M' = A^{1/2}C$ and apply item (2) of Claim A.1 to conclude that $\|A^{1/2}CA^{1/2}\| \leq \|A^{1/2}CB^{1/2}\|$. Set $M'' = CB^{1/2}$ and apply item (1) of the same claim to get that $\|A^{1/2}CB^{1/2}\| \leq \|B^{1/2}CB^{1/2}\|$. It follows that $\|A^{1/2}CA^{1/2}\| \leq \|B^{1/2}CB^{1/2}\|$. The claim follows from the fact that $A^{1/2}CA^{1/2}$ and $B^{1/2}CB^{1/2}$ are symmetric matrices and for all symmetric matrices N , we have that $\|N\| = \rho(N)$. \square

B SMALL SPACE LAPLACIAN SOLVER FOR WEIGHTED GRAPHS

In this section, we show how to compute an approximate pseudoinverse of a Laplacian of a weighted undirected graph deterministically in small space, extending the work of the authors of [43], who only gave an algorithm for undirected multigraphs (graphs with integer edge weights). We sketch the proof of the following lemma.

Lemma 5.4 restated. *Given an undirected connected weighted graph $G = (V, E, w)$ with Laplacian $L = D - A$, and $0 < \varepsilon < 1$, there exists a deterministic algorithm that computes a symmetric PSD matrix \widetilde{L}^+ such that $\widetilde{L}^+ \approx_\varepsilon L^+$, and uses space $O(\log(N \cdot w) \log \log(N \cdot w/\varepsilon))$, where $w = w_{\max}/w_{\min}$ is the ratio of the maximum and minimum edge weights in G and N is the bitlength of the input.*

In [43], the authors assume that multigraphs are given as a list of edges and, hence, the length of the input N is always at least the sum of the edge weights. When moving to general weighted graphs, we will think of rational edge weights as being given as a numerator and a denominator each in binary. This is a more concise representation and, hence, incurs a space complexity that only depends on the weights ratio. One can work with standard floating point representation and get a dependence on the maximal weight.

The following claim about spectral approximation will be useful for proving Lemma 5.4.

CLAIM B.1. *Let A, B, C be symmetric PSD matrices and let $0 < \varepsilon, \varepsilon' < 1$. Then:*

- (1) *If $A \approx_\varepsilon B$, then $A^+ \approx_{\varepsilon/(1-\varepsilon)} B^+$.*
- (2) *If $A \approx_\varepsilon B$ and $B \approx_{\varepsilon'} C$, then $A \approx_{\varepsilon+\varepsilon'+\varepsilon \cdot \varepsilon'} C$.*

See [43] for the proofs.⁹

PROOF SKETCH OF LEMMA 5.4. The algorithm of [43] has two steps, much like most of the efficient Laplacian solvers in the literature. Given the Laplacian L of an undirected multigraph, they first compute a c -spectral approximation to L^+ for some constant $c < 1/2$, using space

⁹The authors of [43] use a variant of spectral approximation that measures closeness with $e^{\pm \varepsilon}$ rather than $(1 \pm \varepsilon)$ but these two notions are equivalent up to constant factors and the proofs of these basic facts are nearly identical for both notions.

$O(\log(N \cdot w_{\max}) \cdot \log \log(N \cdot w_{\max}))$ Then, they boost this to an ε -spectral approximation using an iterative method known as the Richardson iterations. This latter step uses an additional $O(\log(N \cdot w_{\max}) \cdot \log \log \frac{N \cdot w_{\max}}{\varepsilon})$ space, giving the final space complexity. The first step is more delicate and is where the authors required integer edge weights, while Richardson iterations are agnostic to the edge weights and boosts any c -spectral approximation to the pseudoinverse to an ε -spectral approximation in space $O(\log(N \cdot w_{\max}) \cdot \log \log(N \cdot w_{\max})/\varepsilon)$, as long as $c < 1/2$.

Now, given a weighted Laplacian L , we will construct a multigraph having Laplacian L' such that a scalar multiple of $(L')^+$ is a constant spectral approximation to L^+ . Applying the first step of [43] to L' will allow us to compute a constant spectral approximation to $(L')^+$, which, in turn, will be a constant spectral approximation to L^+ by transitivity (Item (2) of Claim B.1). With this, we can boost to an ε -spectral approximation using Richardson iterations.

Set $z = \min\{1, w_{\min}\}$, where w_{\min} is the minimum edge weight in G . Set $\delta = \frac{1}{6}$, $\gamma = \frac{\delta}{1-\delta} = \frac{1}{5}$, and $t = \lceil \log \frac{2n^3}{\delta z} \rceil$. We construct a multigraph G' on n vertices as follows. For each edge $(a, b) \in E(G)$, add an edge between a and b in G' with weight $\lfloor 2^t \cdot w_{ab} \rfloor$. Note that from the way we set t , these new weights are all positive integers. Let $L' = D' - A'$ be the Laplacian of G' and note that for all $a, b \in [n]$, we have that

$$\begin{aligned} |2^{-t} \cdot A'_{ab} - A_{ab}| &= |2^{-t} \cdot \lfloor 2^t \cdot w_{ab} \rfloor - w_{ab}| \\ &\leq 2^{-t}. \end{aligned}$$

Since for all i , $D_{ii} = \sum_j A_{ij}$ and $2^{-t} \cdot D'_{ii} = 2^{-t} \cdot \sum_j A'_{ij}$, it follows that $|2^{-t} \cdot D'_{ii} - D_{ii}| \leq n \cdot 2^{-t}$. Letting $E = 2^{-t} \cdot L' - L$, we can conclude that the sum of absolute values of the entries of each column of E is bounded by $2 \cdot n \cdot 2^{-t}$, denoted by $\|E\|_1 \leq 2 \cdot n \cdot 2^{-t}$.

We will show that $2^{-t} \cdot L' \approx_\delta L$, which, from Claim B.1, will imply that $2^t \cdot (L')^+ = (2^{-t} \cdot L')^+ \approx_\gamma L^+$. To see that $2^{-t} \cdot L' \approx_\delta L$, fix $v \in \mathbb{R}^n$. Since L and L' are both Laplacians of undirected graphs with the same connectivity status, they share a kernel. Thus, for all vectors in the kernel, L and L' have equal quadratic forms and, without loss of generality, we can assume that v is in the orthogonal complement of the kernel.

For any symmetric matrix M , we have that $|v^\top M v| \leq v^\top v \cdot \|M\| \leq v^\top v \cdot \|M\|_1$. Setting $M = E$ and recalling our bound on the $\|\cdot\|_1$ -norm of E , we get that

$$|v^\top E v| \leq v^\top v \cdot 2n \cdot 2^{-t}.$$

From Claim 2.9, we have that $v^\top L v \geq v^\top v \cdot \frac{z}{n^2}$. Combining this with the above gives

$$\begin{aligned} |v^\top E v| &\leq v^\top v \cdot 2n \cdot 2^{-t} \\ &\leq v^\top v \cdot \delta \cdot \frac{z}{n^2} \\ &\leq \delta \cdot v^\top L v, \end{aligned}$$

from which it follows that $2^{-t} \cdot L' \approx_\delta L$. Thus, $2^t \cdot (L')^+ \approx_\gamma L^+$.

Using the first step in [43], we compute \tilde{L} , which is a γ -approximation to $2^t \cdot (L')^+$. Note that the maximum edge weight in G' is upper-bounded by $w_{\max} \cdot 2^t = \text{poly}(n \cdot w_{\max}/w_{\min})$. Thus, the invocation of [43] uses space $O(\log(N \cdot w_{\max}/w_{\min}) \cdot \log \log(N \cdot w_{\max}/w_{\min}))$. From Item (2) of Claim B.1, we have that $\tilde{L} \approx_{2\gamma+\gamma^2} L^+$. Since $2\gamma + \gamma^2 < \frac{1}{2}$, we can apply Richardson iterations to \tilde{L} to compute an ε -spectral approximation to L^+ in total space

$$O(\log(N \cdot w_{\max}/w_{\min}) \log \log(N \cdot w_{\max}/\varepsilon \cdot w_{\min})),$$

as desired. \square

ACKNOWLEDGMENTS

We thank Jelani Nelson for his insights on spectral sparsification via k -wise independent sampling. We also thank Jarosław Błasiok for helpful discussions about random matrices. The first author would like to thank Tselil Schramm and Amnon Ta-Shma for interesting conversations.

REFERENCES

- [1] AmirMahdi Ahmadinejad, Jonathan Kelner, Jack Murtagh, John Peebles, Aaron Sidford, and Salil Vadhan. 2019. High-precision estimation of random walks in small space. *arXiv preprint arXiv:1912.04524* (2019).
- [2] Zeyuan Allen-Zhu, Zhenyu Liao, and Lorenzo Orecchia. 2015. Spectral sparsification and regret minimization beyond matrix multiplicative updates. In *Proceedings of the 47th Annual ACM Symposium on Theory of Computing (STOC'15)*. ACM, 237–245.
- [3] Noga Alon, László Babai, and Alon Itai. 1986. A fast and simple randomized parallel algorithm for the maximal independent set problem. *Journal of Algorithms* 7, 4 (1986), 567–583.
- [4] Noga Alon and Asaf Nussboim. 2008. k -wise independent random graphs. In *Proceedings of the 49th Annual IEEE Symposium on Foundations of Computer Science (FOCS'08)*. IEEE, 813–822.
- [5] Sanjeev Arora and Boaz Barak. 2009. *Computational Complexity: A Modern Approach* (1st ed.). Cambridge University Press, New York, NY, USA.
- [6] Joshua Batson, Daniel A. Spielman, and Nikhil Srivastava. 2014. Twice-Ramanujan sparsifiers. *SIAM Rev.* 56, 2 (2014), 315–334.
- [7] András A. Benczúr and David R. Karger. 1996. Approximating s - t minimum cuts in $\tilde{O}(n^2)$ time. In *STOC*, Vol. 96. Citeseer, 47–55.
- [8] Norman L. Biggs and Alan G. Boshier. 1990. Note on the girth of Ramanujan graphs. *Journal of Combinatorial Theory, Series B* 49, 2 (1990), 190–194.
- [9] Béla Bollobás. 2013. *Modern Graph Theory*. Vol. 184. Springer Science & Business Media.
- [10] Richard Y. Chen, Alex Gittens, and Joel A. Tropp. 2012. The masked sample covariance estimator: An analysis using matrix concentration inequalities. *Information and Inference: A Journal of the IMA* 1, 1 (2012), 2–20.
- [11] L. Paul Chew. 1989. There are planar graphs almost as good as the complete graph. *J. Comput. System Sci.* 39, 2 (1989), 205–219.
- [12] Paul Christiano, Jonathan A. Kelner, Aleksander Madry, Daniel A. Spielman, and Shang-Hua Teng. 2011. Electrical flows, Laplacian systems, and faster approximation of maximum flow in undirected graphs. In *Proceedings of the 43rd Annual ACM Symposium on Theory of Computing (STOC'11)*. ACM, 273–282.
- [13] Timothy Chu, Yu Gao, Richard Peng, Sushant Sachdeva, Saurabh Sawlani, and Junxing Wang. 2018. Graph sparsification, spectral sketches, and faster resistance computation, via short cycle decompositions. In *Proceedings of the 59th Annual IEEE Symposium on Foundations of Computer Science (FOCS'18)*. IEEE, 361–372.
- [14] Michael B. Cohen, Jonathan Kelner, Rasmus Kyng, John Peebles, Richard Peng, Anup B. Rao, and Aaron Sidford. 2018. Solving directed Laplacian systems in nearly-linear time through sparse LU factorizations. In *Proceedings of the 59th Annual IEEE Symposium on Foundations of Computer Science (FOCS'18)*. IEEE, 898–909.
- [15] Michael B. Cohen, Jonathan Kelner, John Peebles, Richard Peng, Anup B. Rao, Aaron Sidford, and Adrian Vladu. 2017. Almost-linear-time algorithms for Markov chains and new spectral primitives for directed graphs. In *Proceedings of the 49th Annual ACM Symposium on Theory of Computing (STOC'17)*. ACM, 410–419.
- [16] Michael B. Cohen, Jonathan Kelner, John Peebles, Richard Peng, Aaron Sidford, and Adrian Vladu. 2016. Faster algorithms for computing the stationary distribution, simulating random walks, and more. In *Proceedings of the 57th Annual IEEE Symposium on Foundations of Computer Science (FOCS'16)*. IEEE, 583–592.
- [17] Michael B. Cohen, Rasmus Kyng, Gary L. Miller, Jakub W. Pachocki, Richard Peng, Anup B. Rao, and Shen Chen Xu. 2014. Solving SDD linear systems in nearly $m \log^{1/2} n$ time. In *Proceedings of the 46th Annual ACM Symposium on Theory of Computing (STOC'14)*. ACM, 343–352.
- [18] Michael B Cohen, Aleksander Madry, Piotr Sankowski, and Adrian Vladu. 2017. Negative-weight shortest paths and unit capacity minimum cost flow in $\tilde{O}(m^{10/7} \log W)$ time. In *Proceedings of the 28th Annual ACM-SIAM Symposium on Discrete Algorithms (SODA'17)*. SIAM, 752–771.
- [19] Dean Doron, François Le Gall, and Amnon Ta-Shma. 2017. Probabilistic logarithmic-space algorithms for Laplacian solvers. In *Approximation, Randomization, and Combinatorial Optimization. Algorithms and Techniques (APPROX/RANDOM'17)*. Schloss Dagstuhl-Leibniz-Zentrum fuer Informatik, 41:1–41:20.
- [20] Paul Erdős and Horst Sachs. 1963. Reguläre graphen gegebener taillenweite mit minimaler knotenzahl. *Wiss. Z. Martin-Luther-Univ. Halle-Wittenberg Math.-Natur. Reihe* 12, 251–257 (1963), 22.
- [21] Geoffrey Exoo and Robert Jajcay. 2008. Dynamic cage survey. *Electron. J. Combin* 15, 16 (2008), 4.

- [22] Ronald M. Foster. 1949. The average impedance of an electrical network. *Contributions to Applied Mechanics (Reissner Anniversary Volume)* (1949), 333–340.
- [23] Oded Goldreich. 2008. Computational complexity: A conceptual perspective. *ACM SIGACT News* 39, 3 (2008), 35–39.
- [24] A. Joffe. 1974. On a set of almost deterministic k -independent random variables. *The Annals of Probability* 2, 1 (1974), 161–162.
- [25] Daniel Kane, Raghu Meka, and Jelani Nelson. 2011. Almost optimal explicit Johnson-Lindenstrauss families. In *Approximation, Randomization, and Combinatorial Optimization. Algorithms and Techniques (APPROX/RANDOM'11)*. Springer, 628–639.
- [26] Jonathan A. Kelner, Yin Tat Lee, Lorenzo Orecchia, and Aaron Sidford. 2014. An almost-linear-time algorithm for approximate max flow in undirected graphs, and its multicommodity generalizations. In *Proceedings of the 25th Annual ACM-SIAM Symposium on Discrete Algorithms (SODA'14)*. SIAM, 217–226.
- [27] Jonathan A. Kelner and Aleksander Madry. 2009. Faster generation of random spanning trees. In *Proceedings of the 50th Annual IEEE Symposium on Foundations of Computer Science (FOCS'19)*. IEEE, 13–21.
- [28] Jonathan A. Kelner, Lorenzo Orecchia, Aaron Sidford, and Zeyuan Allen Zhu. 2013. A simple, combinatorial algorithm for solving SDD systems in nearly-linear time. In *Proceedings of the 45th Annual ACM Symposium on Theory of Computing (STOC'13)*. ACM, 911–920.
- [29] Ioannis Koutis, Gary L. Miller, and Richard Peng. 2011. A nearly- $m \log n$ time solver for SDD linear systems. In *Proceedings of the 52nd Annual IEEE Symposium on Foundations of Computer Science (FOCS'11)*. IEEE, 590–598.
- [30] Ioannis Koutis, Gary L. Miller, and Richard Peng. 2014. Approaching optimality for solving SDD linear systems. *SIAM J. Comput.* 43, 1 (2014), 337–354.
- [31] Ioannis Koutis, Gary L. Miller, and David Tolliver. 2009. Combinatorial preconditioners and multilevel solvers for problems in computer vision and image processing. In *International Symposium on Visual Computing*. Springer, 1067–1078.
- [32] Rasmus Kyng, Yin Tat Lee, Richard Peng, Sushant Sachdeva, and Daniel A. Spielman. 2016. Sparsified Cholesky and multigrid solvers for connection Laplacians. In *Proceedings of the 48th Annual ACM Symposium on Theory of Computing (STOC'16)*. ACM, 842–850.
- [33] Rasmus Kyng and Sushant Sachdeva. 2016. Approximate Gaussian elimination for Laplacians—fast, sparse, and simple. In *Proceedings of the 57th Annual IEEE Symposium on Foundations of Computer Science (FOCS'16)*. IEEE, 573–582.
- [34] Felix Lazebnik, Vasily A. Ustimenko, and Andrew J. Woldar. 1995. A new series of dense graphs of high girth. *Bulletin of the American Mathematical Society* 32, 1 (1995), 73–79.
- [35] Yin Tat Lee, Satish Rao, and Nikhil Srivastava. 2013. A new approach to computing maximum flows using electrical flows. In *Proceedings of the 45th Annual ACM Symposium on Theory of Computing (STOC'13)*. ACM, 755–764.
- [36] Yin Tat Lee and Aaron Sidford. 2013. Efficient accelerated coordinate descent methods and faster algorithms for solving linear systems. In *Proceedings of the 54th Annual IEEE Symposium on Foundations of Computer Science (FOCS'13)*. IEEE, 147–156.
- [37] Yin Tat Lee and He Sun. 2015. Constructing linear-sized spectral sparsification in almost-linear time. In *Proceedings of the 56th IEEE Annual Symposium on Foundations of Computer Science (FOCS'15)*. IEEE, 250–269.
- [38] Yin Tat Lee and He Sun. 2017. An SDP-based algorithm for linear-sized spectral sparsification. In *Proceedings of the 49th Annual ACM Symposium on Theory of Computing (STOC'17)*. ACM, 678–687.
- [39] Yang P. Liu and Aaron Sidford. 2019. Faster energy maximization for faster maximum flow. *arXiv preprint arXiv:1910.14276* (2019).
- [40] Alexander Lubotzky, Ralph Phillips, and Peter Sarnak. 1988. Ramanujan graphs. *Combinatorica* 8, 3 (1988), 261–277.
- [41] Michael Luby. 1986. A simple parallel algorithm for the maximal independent set problem. *SIAM Journal on Computing* 15, 4 (1986), 1036–1053.
- [42] Aleksander Madry, Damian Straszak, and Jakub Tarnawski. 2014. Fast generation of random spanning trees and the effective resistance metric. In *Proceedings of the 26th Annual ACM-SIAM Symposium on Discrete Algorithms (SODA'14)*. SIAM, 2019–2036.
- [43] Jack Murtagh, Omer Reingold, Aaron Sidford, and Salil P. Vadhan. 2017. Derandomization beyond connectivity: Undirected Laplacian systems in nearly logarithmic space. In *Proceedings of the 58th IEEE Annual Symposium on Foundations of Computer Science (FOCS'17)*. IEEE, 801–812.
- [44] Jack Murtagh, Omer Reingold, Aaron Sidford, and Salil P. Vadhan. 2019. Deterministic approximation of random walks in small space. In *Approximation, Randomization, and Combinatorial Optimization. Algorithms and Techniques (APPROX/RANDOM'19)*. Schloss Dagstuhl-Leibniz-Zentrum fuer Informatik, 42:1–42:22.
- [45] Lorenzo Orecchia, Sushant Sachdeva, and Nisheeth K. Vishnoi. 2012. Approximating the exponential, the Lanczos method and an $\tilde{O}(m)$ -time spectral algorithm for balanced separator. In *Proceedings of the 44th Annual ACM Symposium on Theory of Computing (STOC'12)*. ACM, 1141–1160.

- [46] Richard Peng and Daniel A. Spielman. 2014. An efficient parallel solver for SDD linear systems. In *Proceedings of the 46th Annual ACM Symposium on Theory of Computing (STOC'14)*. ACM, 333–342.
- [47] Robert Qiu and Michael Wicks. 2014. *Cognitive Networked Sensing and Big Data*. Springer.
- [48] Omer Reingold. 2008. Undirected connectivity in log-space. *J. ACM* 55, 4 (2008).
- [49] Eyal Rozenman and Salil Vadhan. 2005. Derandomized squaring of graphs. In *Approximation, Randomization and Combinatorial Optimization. Algorithms and Techniques*. Springer, 436–447.
- [50] Mark Rudelson and Roman Vershynin. 2007. Sampling from large matrices: An approach through geometric functional analysis. *Journal of the ACM (JACM)* 54, 4 (2007), 21.
- [51] Aaron Schild. 2018. An almost-linear time algorithm for uniform random spanning tree generation. In *Proceedings of the 50th Annual ACM Symposium on Theory of Computing (STOC'18)*. ACM, 214–227.
- [52] Daniel A. Spielman. 2015. Yale Applied Mathematics 561/Computer Science 662, Lecture Notes: Spectral Graph Theory. <http://www.cs.yale.edu/homes/spielman/561/lect17-15.pdf>
- [53] Daniel A. Spielman and Nikhil Srivastava. 2011. Graph sparsification by effective resistances. *SIAM Journal on Computing* 40, 6 (2011), 1913–1926.
- [54] Daniel A. Spielman and Shang-Hua Teng. 2004. Nearly-linear time algorithms for graph partitioning, graph sparsification, and solving linear systems. In *Proceedings of the 36th Annual ACM Symposium on Theory of Computing (STOC'04)*. 81–90.
- [55] Anastasios Zouzias. 2012. A matrix hyperbolic cosine algorithm and applications. In *International Colloquium on Automata, Languages, and Programming (ICALP'12)*. Springer, 846–858.

Received 6 October 2021; revised 4 December 2023; accepted 7 December 2023